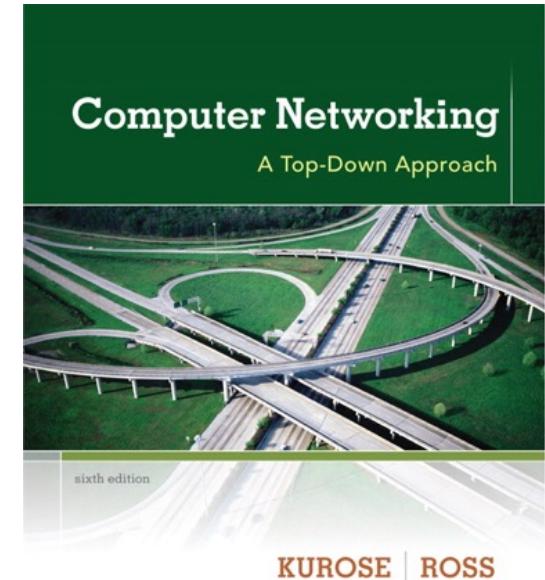


Chapter 2

Application Layer



*Computer
Networking: A Top
Down Approach*
6th edition
Jim Kurose, Keith Ross
Addison-Wesley
March 2012

©

All material copyright 1996-2012
J.F Kurose and K.W. Ross, All Rights Reserved

Modified by Prof. Guoliang Xing, Prof. Man Hon
Cheung, Prof. Jack Y. B. Lee

Application Layer 2-1

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Socket programming with UDP and TCP

Chapter 2: application layer

our goals:

- ❖ conceptual, implementation aspects of network application protocols
 - transport-layer service models
 - client-server paradigm
 - peer-to-peer paradigm
- ❖ learn about protocols by examining popular application-level protocols
 - HTTP
 - FTP
 - SMTP / POP3 / IMAP
 - DNS
- ❖ creating network applications
 - socket API

Some network apps

- ❖ e-mail
- ❖ web
- ❖ text messaging
- ❖ remote login
- ❖ P2P file sharing
- ❖ multi-user network games
- ❖ streaming stored video
(YouTube, Hulu, Netflix)
- ❖ voice over IP (e.g., Skype)
- ❖ real-time video conferencing
- ❖ social networking
- ❖ search
- ❖ ...
- ❖ ...

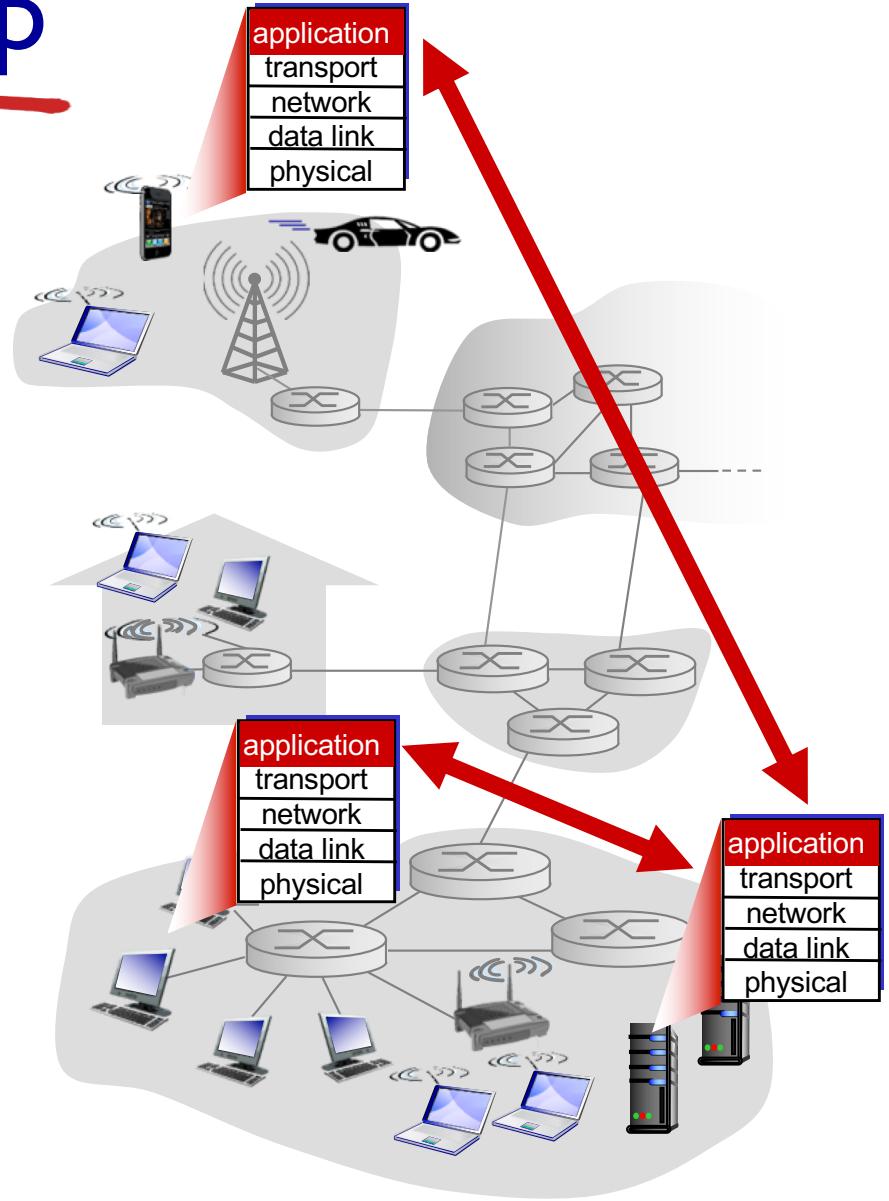
Creating a network app

write programs that:

- ❖ run on (different) end systems
- ❖ communicate over network
- ❖ e.g., web server software communicates with browser software

no need to write software for network-core devices

- ❖ network-core devices do not run user applications
- ❖ applications on end systems allows for rapid app development, propagation

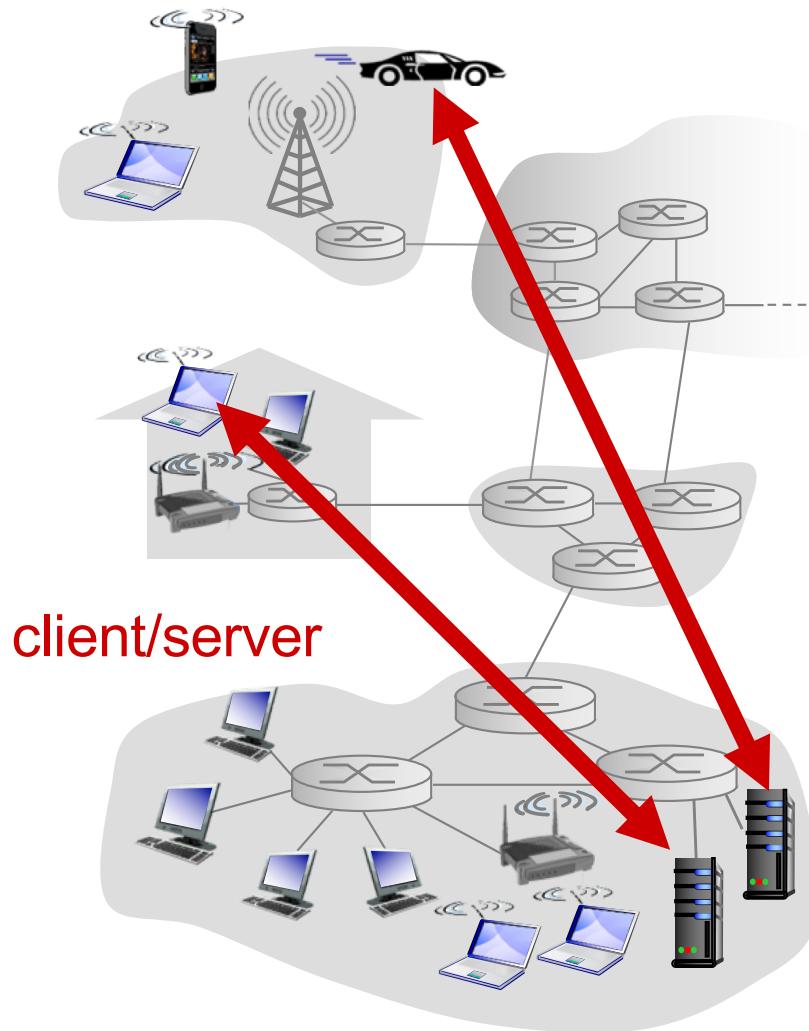


Application architectures

possible structure of applications:

- ❖ client-server
- ❖ peer-to-peer (P2P)

Client-server architecture



server:

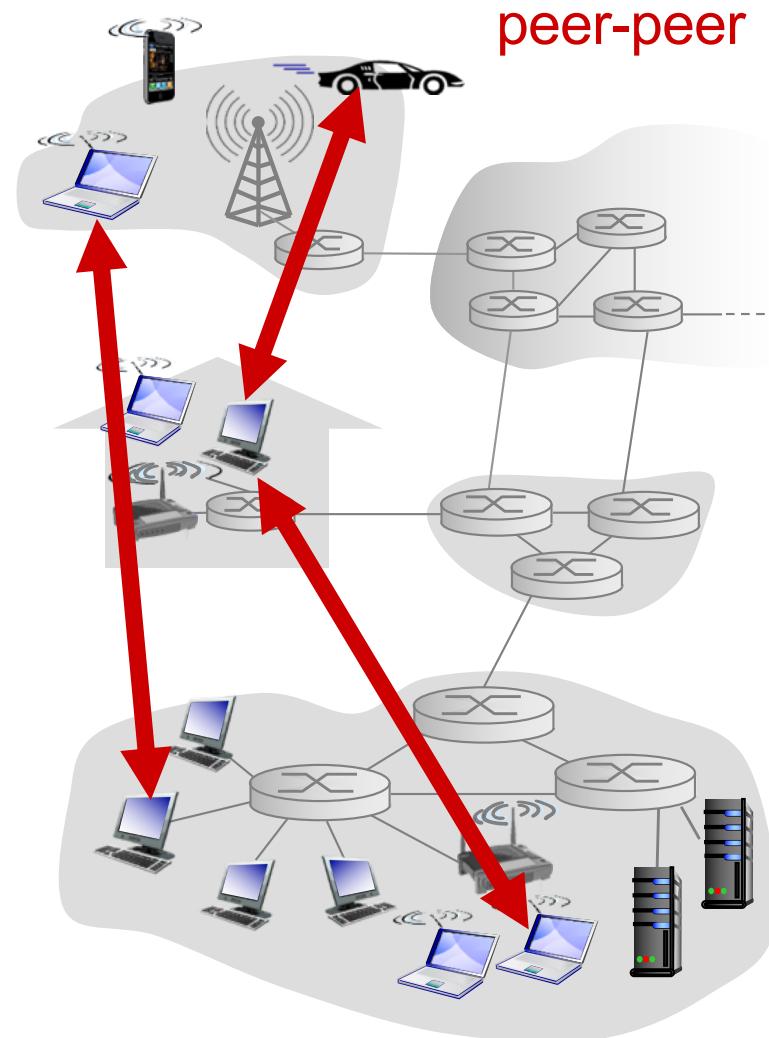
- ❖ always-on host
- ❖ permanent IP address
- ❖ **data centers** for scaling

clients:

- ❖ communicate with server
- ❖ may be intermittently connected
- ❖ may have dynamic IP addresses
- ❖ do not communicate directly with each other

P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- ❖ peers are intermittently connected and change IP addresses
 - complex management



Processes communicating

process: program running within a host

Process communications:

- ❖ within **same host**, two processes communicate using inter-process communication defined by operating system
- ❖ processes in **different hosts** communicate by exchanging **messages**

clients, servers

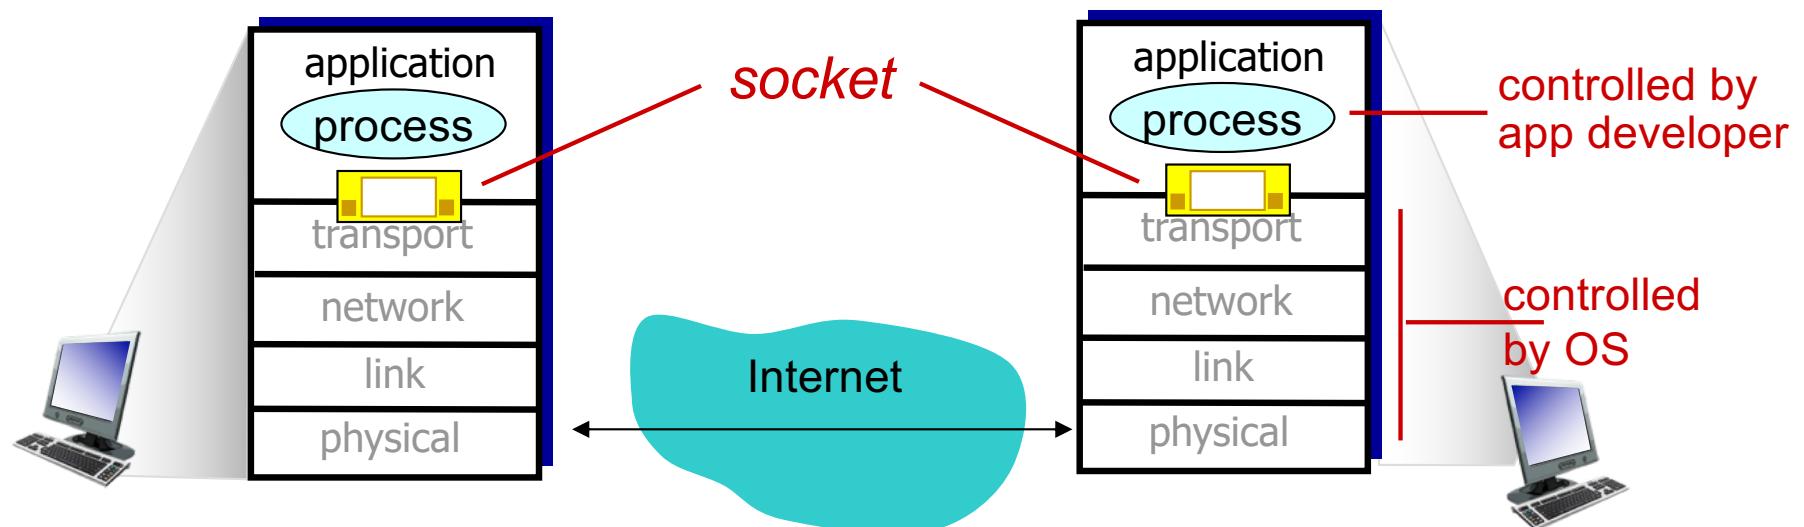
client process: process that initiates communication

server process: process that waits to be contacted

- ❖ aside: applications with P2P architectures have client processes & server processes

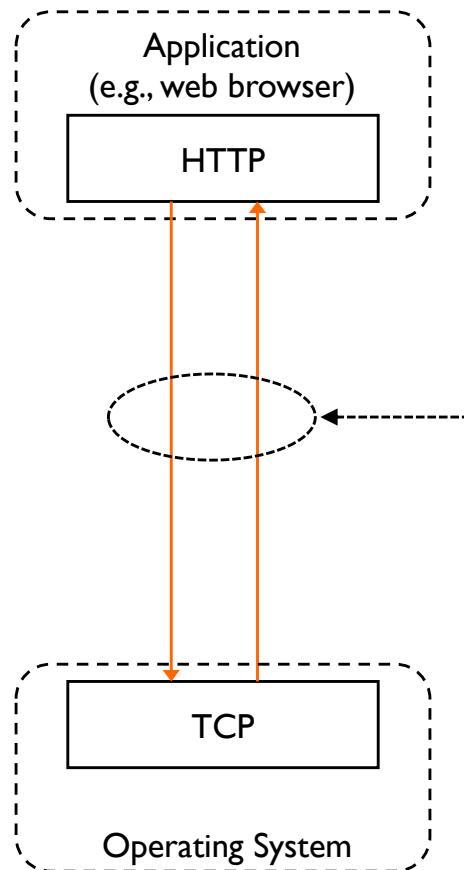
Sockets

- ❖ process sends/receives messages to/from the network through a **software interface** called a **socket**
- ❖ socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Sockets

- ❖ Accessing transport services via sockets API:



Network Application Programming Interface (API) (e.g., Sockets)

```
// Prepare address structures //
TCP_PeerAddr      = new sockaddr_in;
memset (TCP_PeerAddr, 0, sizeof(struct sockaddr_in));
TCP_PeerAddr->sin_family          = AF_INET;
TCP_PeerAddr->sin_port            = htons(0x4747);
TCP_PeerAddr->sin_addr.s_addr     = inet_addr("137.189.97.120");

// Create a socket and accept for incoming connection //
SOCKET Sockfd = socket(AF_INET, SOCK_STREAM, 0);
int retval = connect(Sockfd, (struct sockaddr *)TCP_PeerAddr,
                     sizeof(struct sockaddr_in));
if (retval == SOCKET_ERROR)
    // connect error
else // (retval == 0)
    // connect successful ...
```

Addressing processes

- ❖ to receive messages, process must have *identifier*
- ❖ host device has unique 32-bit IP address
- ❖ Q: does IP address of host on which process runs suffice for identifying the process?
 - A: no, *many processes can be running on same host*
- ❖ *identifier* includes both IP address and port numbers associated with process on host.
- ❖ example port numbers:
 - HTTP server: 80
 - mail server: 25
- ❖ to send HTTP message to gaia.cs.umass.edu web server:
 - IP address: 128.119.245.12
 - port number: 80
- ❖ more shortly...

App-layer protocol defines

1. types of messages exchanged,
 - e.g., request, response
2. message syntax:
 - what fields in messages & how fields are delineated
3. message semantics
 - meaning of information in fields
4. rules for when and how processes send & respond to messages

Two main types of app-layer protocols:

1. open protocols:
 - ❖ defined in RFCs
 - ❖ allows for interoperability
 - ❖ e.g., HTTP, SMTP
2. proprietary protocols:
 - ❖ Not defined in RFCs or elsewhere
 - ❖ A single development team has complete control over what goes in the code
 - ❖ e.g., Skype

What transport service does an app need?

data integrity

- ❖ some apps (e.g., file transfer, web transactions) require **100% reliable** data transfer
- ❖ other apps (e.g., audio) can **tolerate some loss**

timing

- ❖ some apps (e.g., Internet telephony, interactive games) require low **delay** to be “effective”

throughput

- ❖ some apps (e.g., multimedia) require **minimum amount of throughput** to be “effective”
- ❖ other apps (“**elastic apps**”) make use of whatever throughput they get

security

- ❖ encryption, data integrity,

...

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	
interactive games	loss-tolerant	few kbps up	yes, few secs
text messaging	no loss	elastic	yes, 100' s msec yes and no

Internet transport protocols services

TCP service:

- ❖ ***reliable transport*** between sending and receiving process
- ❖ ***flow control***: sender won't overwhelm **receiver**
- ❖ ***congestion control***: throttle sender when **network** overloaded
- ❖ ***does not provide***: timing, minimum throughput guarantee, security
- ❖ ***connection-oriented***: setup required between client and server processes

UDP service:

- ❖ ***unreliable data transfer*** between sending and receiving process
- ❖ ***does not provide***: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

Internet apps: application, transport protocols

application	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

Securing TCP

Problem of TCP & UDP:

no encryption

- ❖ cleartext passwords sent into socket traverse Internet in cleartext

Solution: **Secure Sockets Layer (SSL)**

- ❖ provides encrypted TCP connection
- ❖ data integrity
- ❖ end-point authentication

SSL is at app layer

- ❖ Apps use SSL libraries, which “talk” to TCP

SSL socket API

- ❖ cleartext passwords sent into socket traverse Internet encrypted
- ❖ See Chapter 7

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Socket programming with UDP and TCP

Web and HTTP

First, a review...

- ❖ web page consists of *base HTML-file* which includes *several referenced objects*
 - object can be HTML file, JPEG image, Java applet, audio file,...
 - each object is addressable by a *URL*, e.g.,

www.someschool.edu/someDept/pic.gif

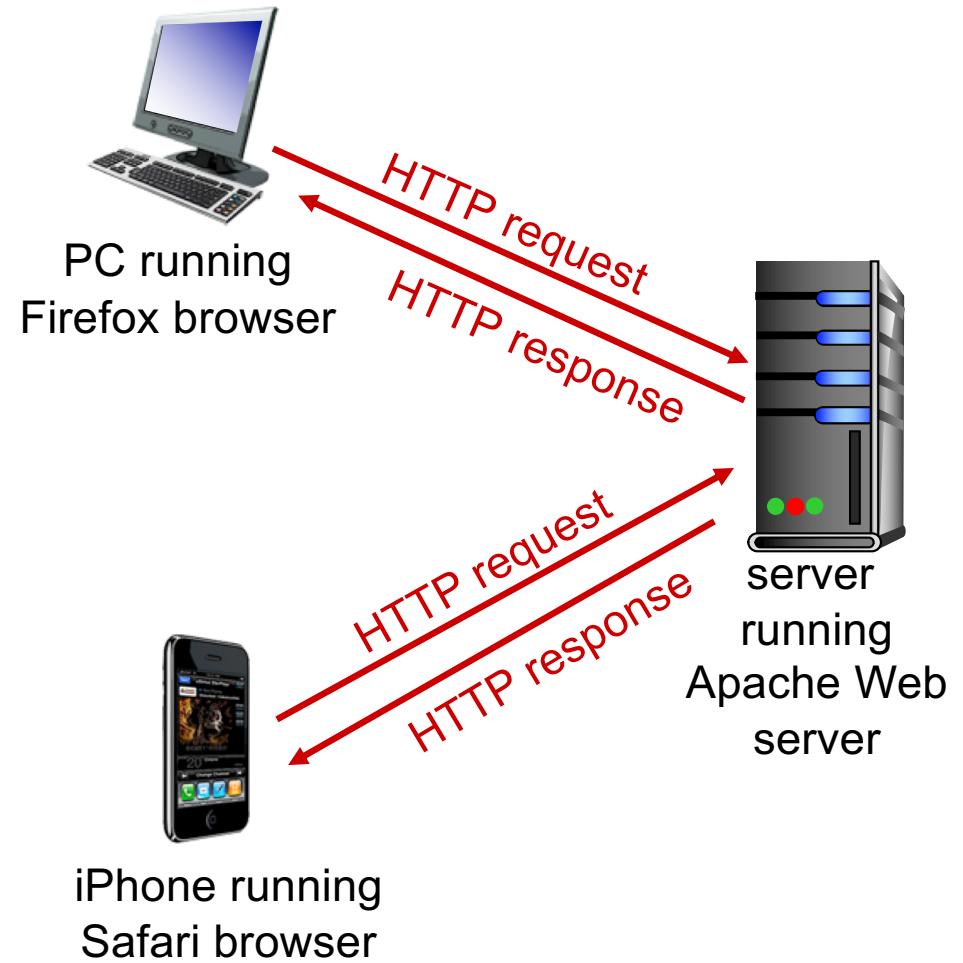
host name

path name

HTTP overview

HTTP: hypertext transfer protocol

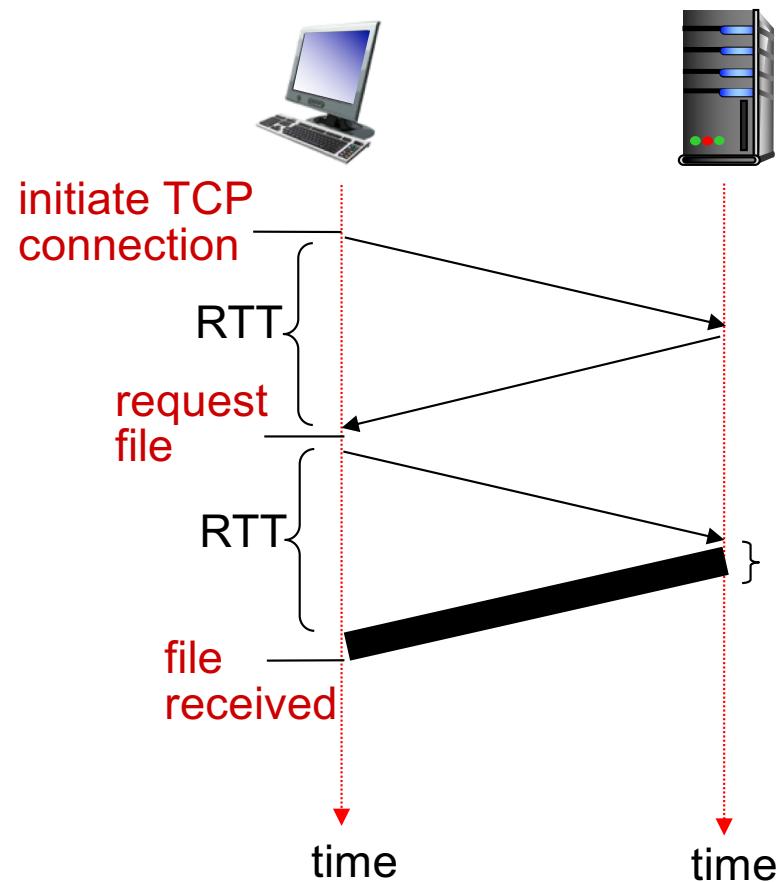
- ❖ Web's application layer protocol
- ❖ client/server model
 - client: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - server: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

uses TCP:

1. client initiates **TCP connection** (creates socket) to server, **port 80**
2. server accepts TCP connection from client
3. HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
4. TCP connection closed



HTTP overview (continued)

*HTTP is “**stateless**”*

- ❖ server maintains no information about past client requests

protocols that maintain “state” are complex!

- ❖ past **history (state)** must be maintained
- ❖ if server/client crashes, their views of “state” may be inconsistent, must be reconciled

aside

HTTP connections

Question: should each request/response pair be sent over a **separate or the same TCP connection?**

non-persistent HTTP

- ❖ at most one object sent over TCP connection
 - connection then closed
- ❖ downloading multiple objects required multiple connections

persistent HTTP

- ❖ multiple objects can be sent over single TCP connection between client, server

Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,
references to 10
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time
↓

Non-persistent HTTP (cont.)

time
↓

5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects



4. HTTP server closes TCP connection.

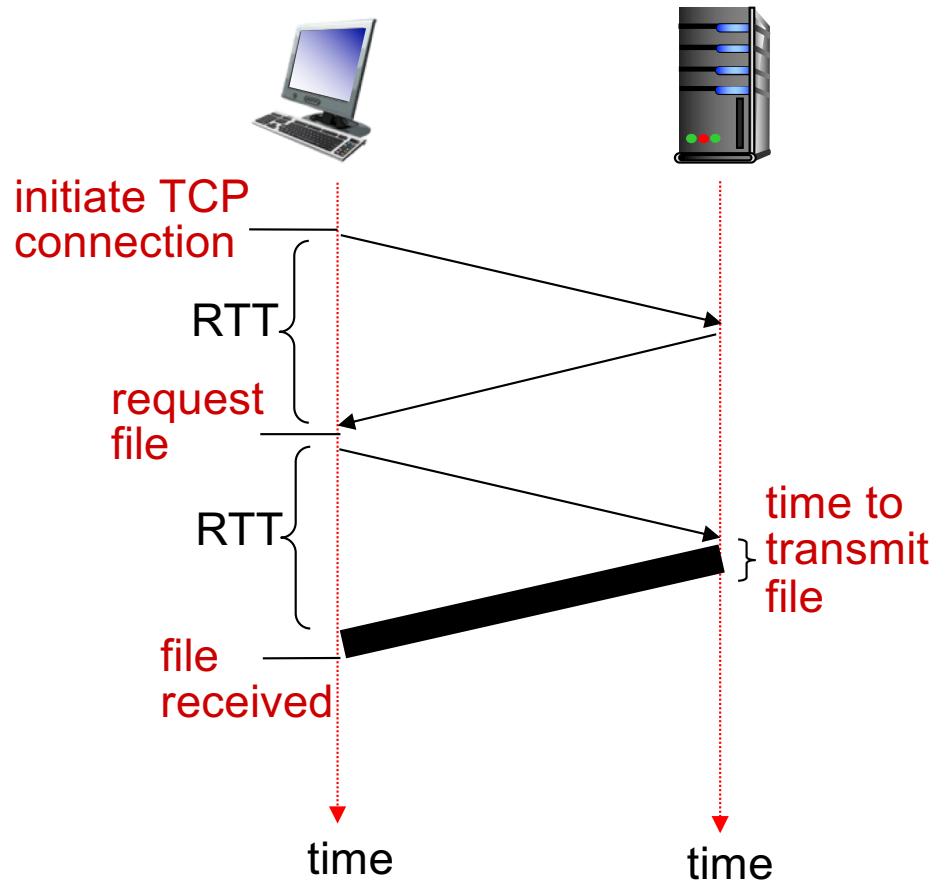
6. Steps 1-5 repeated for each of 10 jpeg objects

Non-persistent HTTP: response time

Round-trip time (RTT): time for a small packet to travel from client to server and back

HTTP response time:

- ❖ one RTT to **initiate TCP connection**
- ❖ one RTT for **HTTP request** and first few bytes of **HTTP response** to return
- ❖ **file transmission** time
- ❖ non-persistent HTTP response time =
$$2\text{RTT} + \text{file transmission time}$$

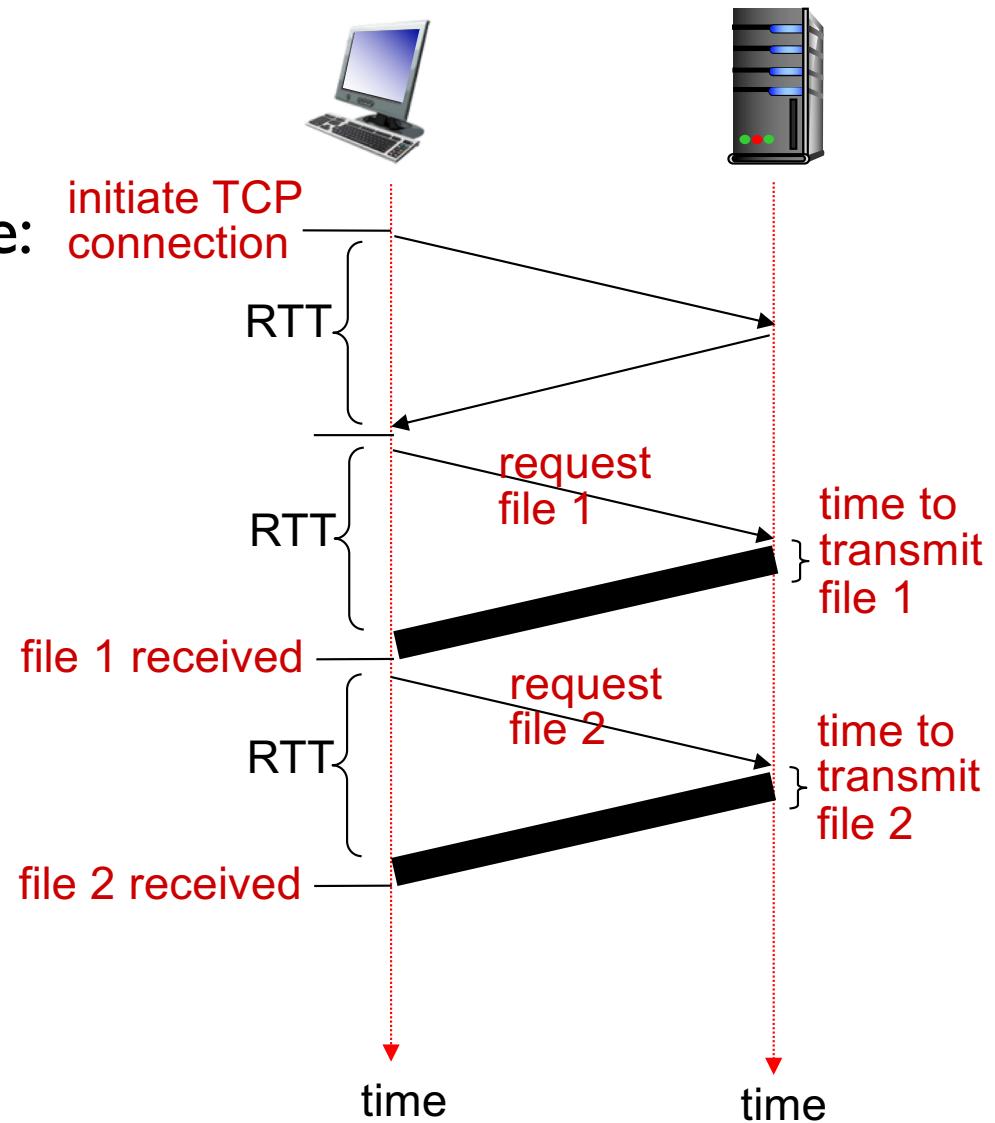


Persistent HTTP: response time

First request is same as non-persistent HTTP

Subsequent HTTP response time:

- ❖ one RTT for HTTP request and first few bytes of HTTP response to return
- ❖ file transmission time
- ❖ persistent HTTP response time =
$$\text{RTT} + \text{file transmission time}$$



Non-Persistent & Persistent HTTP

non-persistent HTTP issues:

- ❖ requires 2 RTTs per object
- ❖ OS overhead for each TCP connection
- ❖ browsers often open **parallel TCP connections** to fetch referenced objects

persistent HTTP:

- ❖ server leaves connection open after sending response
- ❖ subsequent HTTP messages between same client/server sent over open connection
- ❖ client sends requests as soon as it encounters a referenced object

Concurrent HTTP

- ❖ What?
 - Not part of the HTTP specification
 - Implemented by almost all web browsers
 - After the initial HTTP connection which retrieves the HTML body, initiate multiple (4~6) HTTP sessions (per domain) to retrieve multiple objects (e.g., images) in parallel.
- ❖ Why?
 - Reduce delay in waiting for previous HTTP response to complete transfer (aka head-of-line blocking)
 - Potentially overlaps the server-side processing time of multiple HTTP requests
 - Effectively multiplies the TCP congestion window size by the number of TCP connections, resulting in higher aggregate throughput
- ❖ Shortcomings:
 - Use up more resources at server
 - Fairness to competing data flows

User-server state: cookies

Q: As HTTP server is stateless, how can a Web site identify users?

A: **Cookies**, defined in [RFC 6265], allow sites to keep track of users.

example:

- ❖ Susan always access Internet from PC
- ❖ visits specific **e-commerce site** for first time
- ❖ when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

Cookies (continued)

*what cookies can be used
for:*

- ❖ authorization
- ❖ shopping carts
- ❖ recommendations
- ❖ user session state (Web e-mail)

aside
cookies and privacy:

- ❖ cookies permit sites to learn a lot about you
- ❖ you may supply name and e-mail to sites

how to keep “state”:

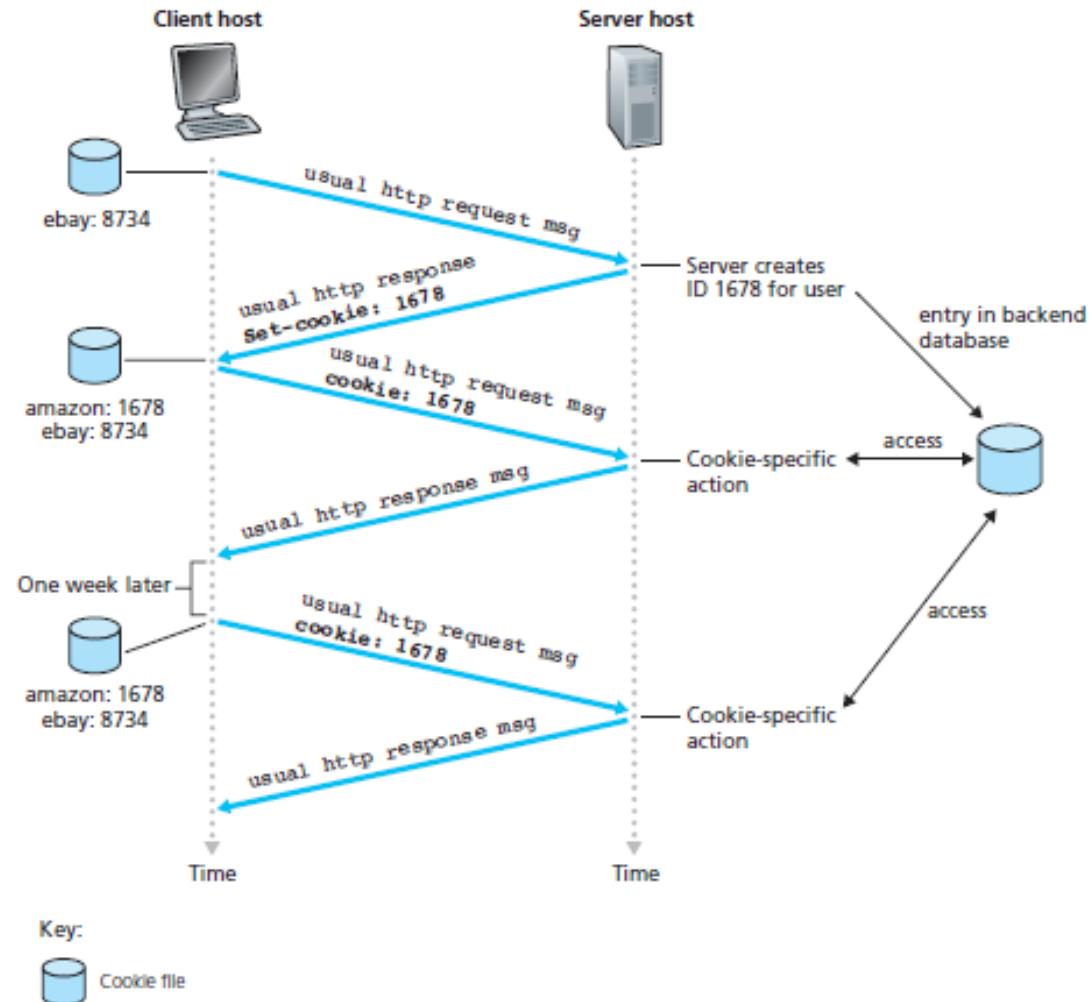
- ❖ protocol endpoints: maintain state at sender/receiver over multiple transactions
- ❖ cookies: http messages carry state

User-server state: cookies

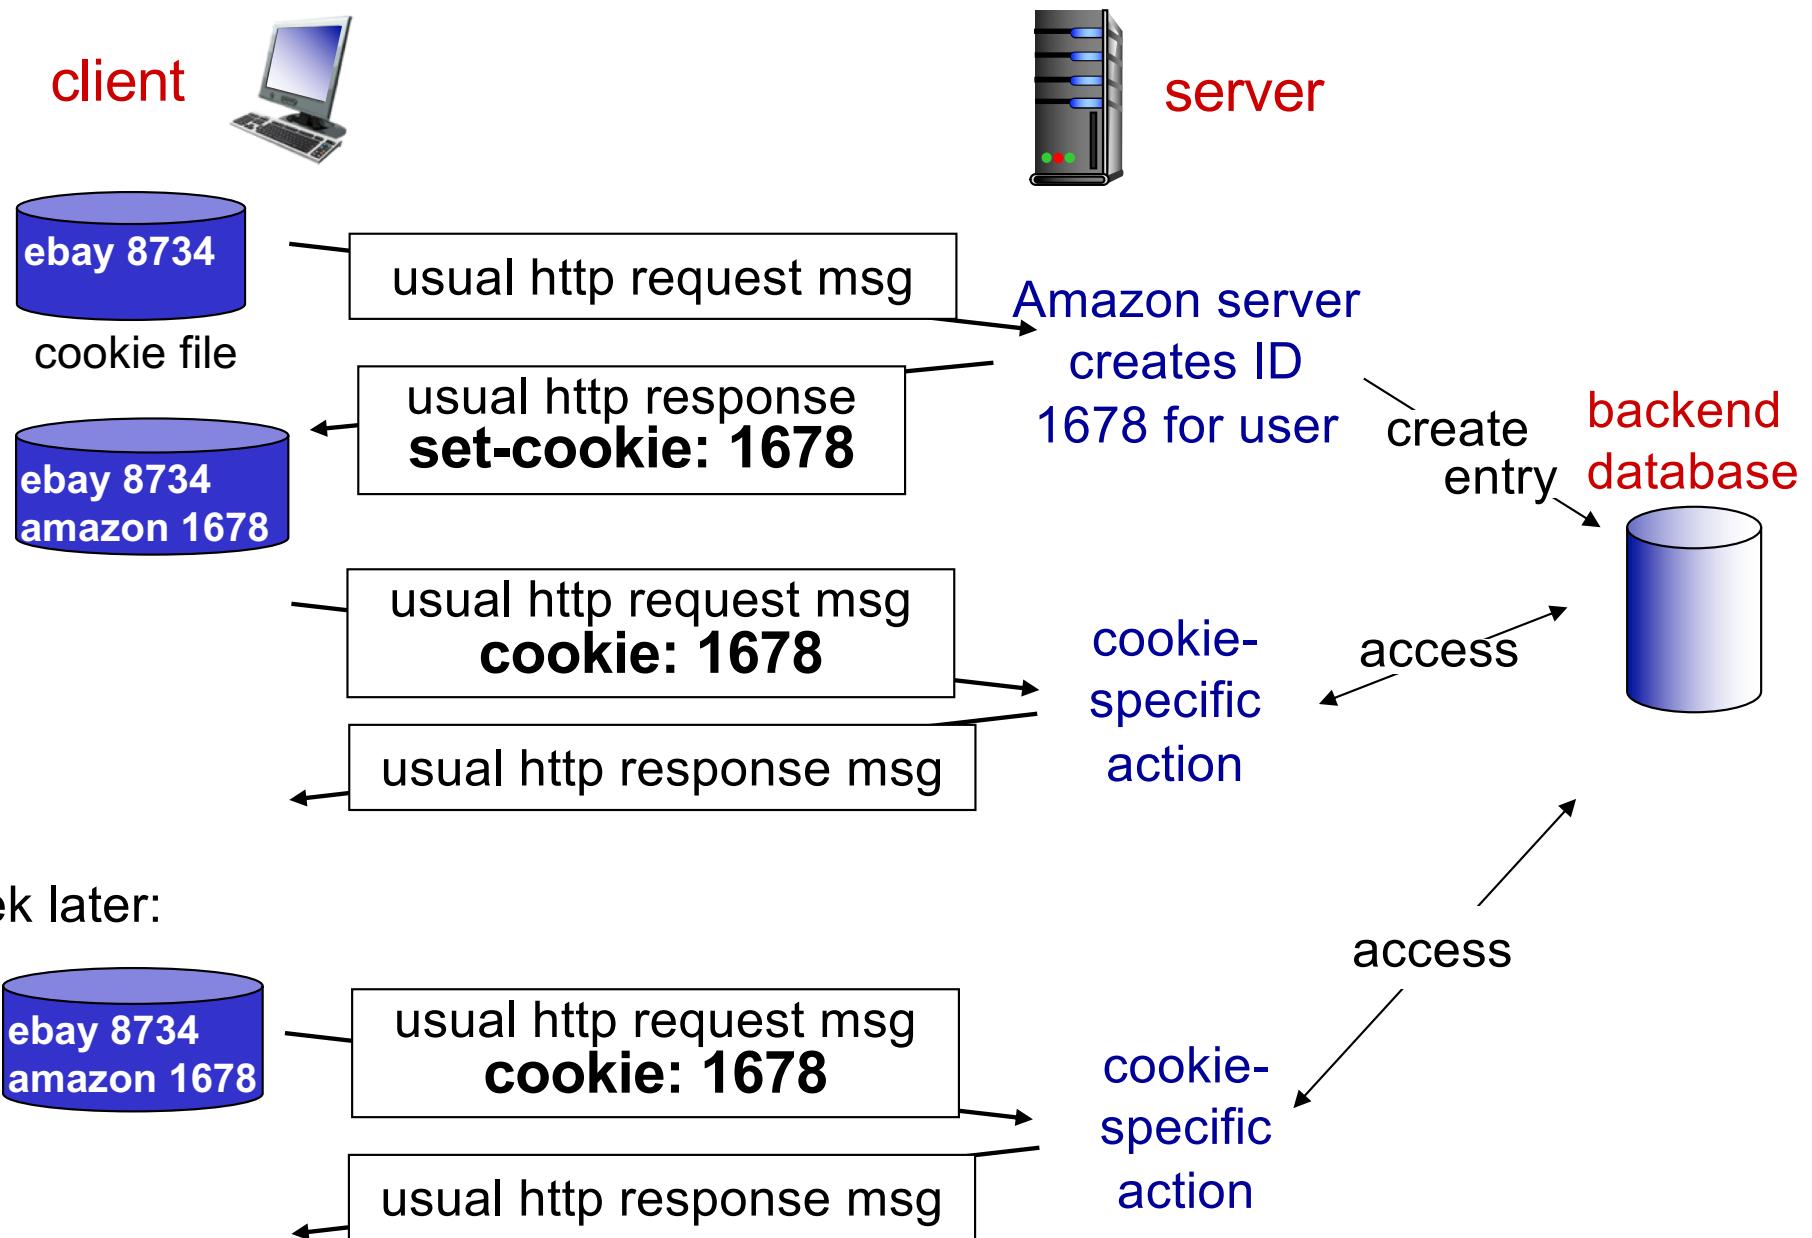
many Web sites use cookies

four components:

- 1) cookie header line of **HTTP response message**
- 2) cookie header line in next **HTTP request message**
- 3) **cookie file** kept on user's host, managed by user's browser
- 4) **backend database** at Web site



Cookies: keeping “state” (cont.)



Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

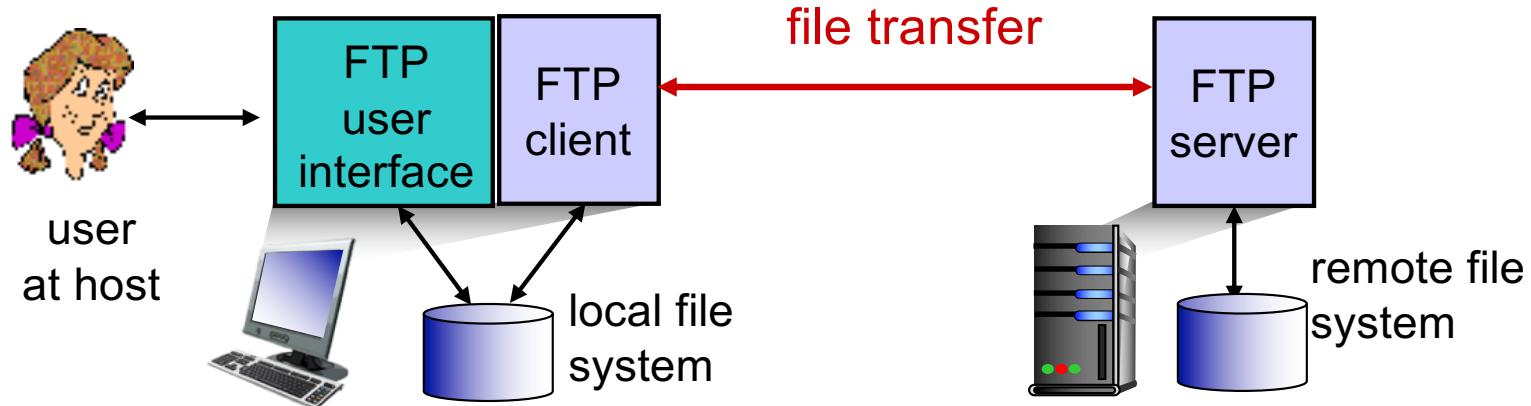
- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Socket programming with UDP and TCP

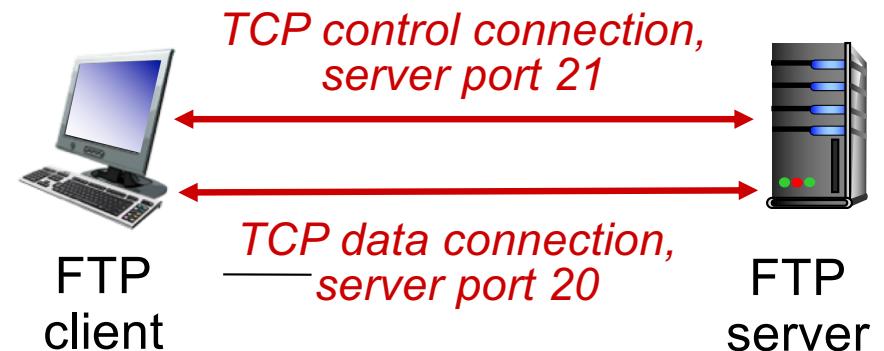
FTP: the file transfer protocol



- ❖ transfer file to/from remote host
- ❖ **client/server model**
 - *client*: side that initiates transfer (either to/from remote)
 - *server*: remote host
- ❖ **ftp server: port 21**

FTP: separate control, data connections

- ❖ FTP client contacts FTP server at port 21, using TCP
- ❖ Control connection (port 21):
 - authentication
 - client browses remote directory, sends commands
- ❖ Data connection (port 20):
 - when server receives file transfer command, **server** opens 2nd TCP data connection (for file) to client
 - after transferring one file, server closes data connection



- server opens another TCP data connection to transfer another file
- ❖ FTP server maintains “state”: current directory, earlier authentication

FTP commands, responses

sample commands:

- ❖ sent as ASCII text over control channel
- ❖ **USER *username***
- ❖ **PASS *password***
- ❖ **LIST** return list of file in current directory
- ❖ **RETR *filename*** retrieves (gets) file
- ❖ **STOR *filename*** stores (puts) file onto remote host

sample return codes

- ❖ status code and phrase (as in HTTP)
- ❖ **331 Username OK, password required**
- ❖ **125 data connection already open; transfer starting**
- ❖ **425 Can't open data connection**
- ❖ **452 Error writing file**

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Socket programming with UDP and TCP

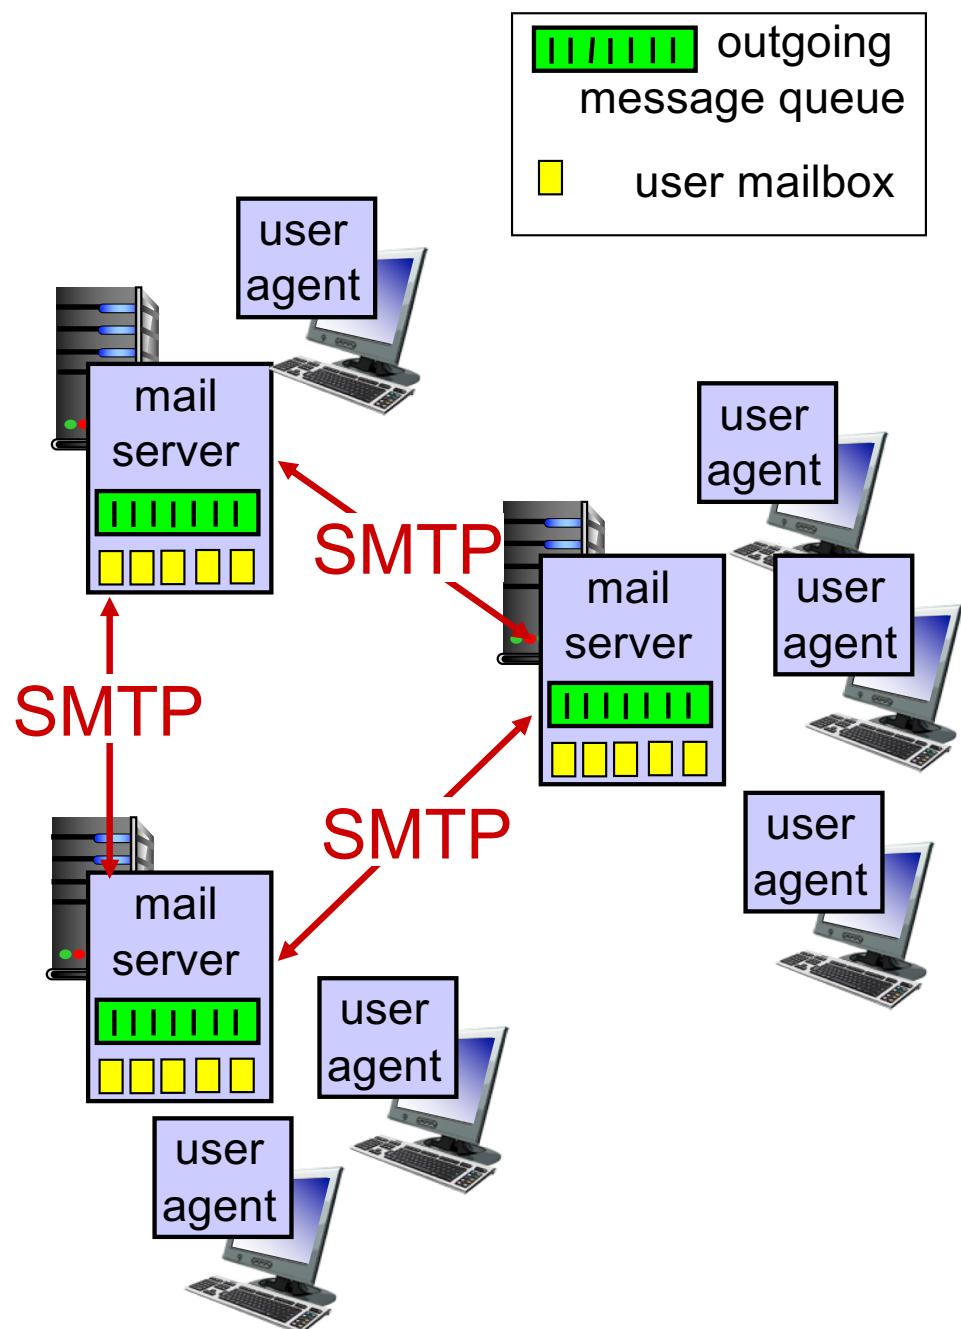
Electronic mail

Three major components:

1. user agents
2. mail servers
3. simple mail transfer protocol: SMTP

User Agent

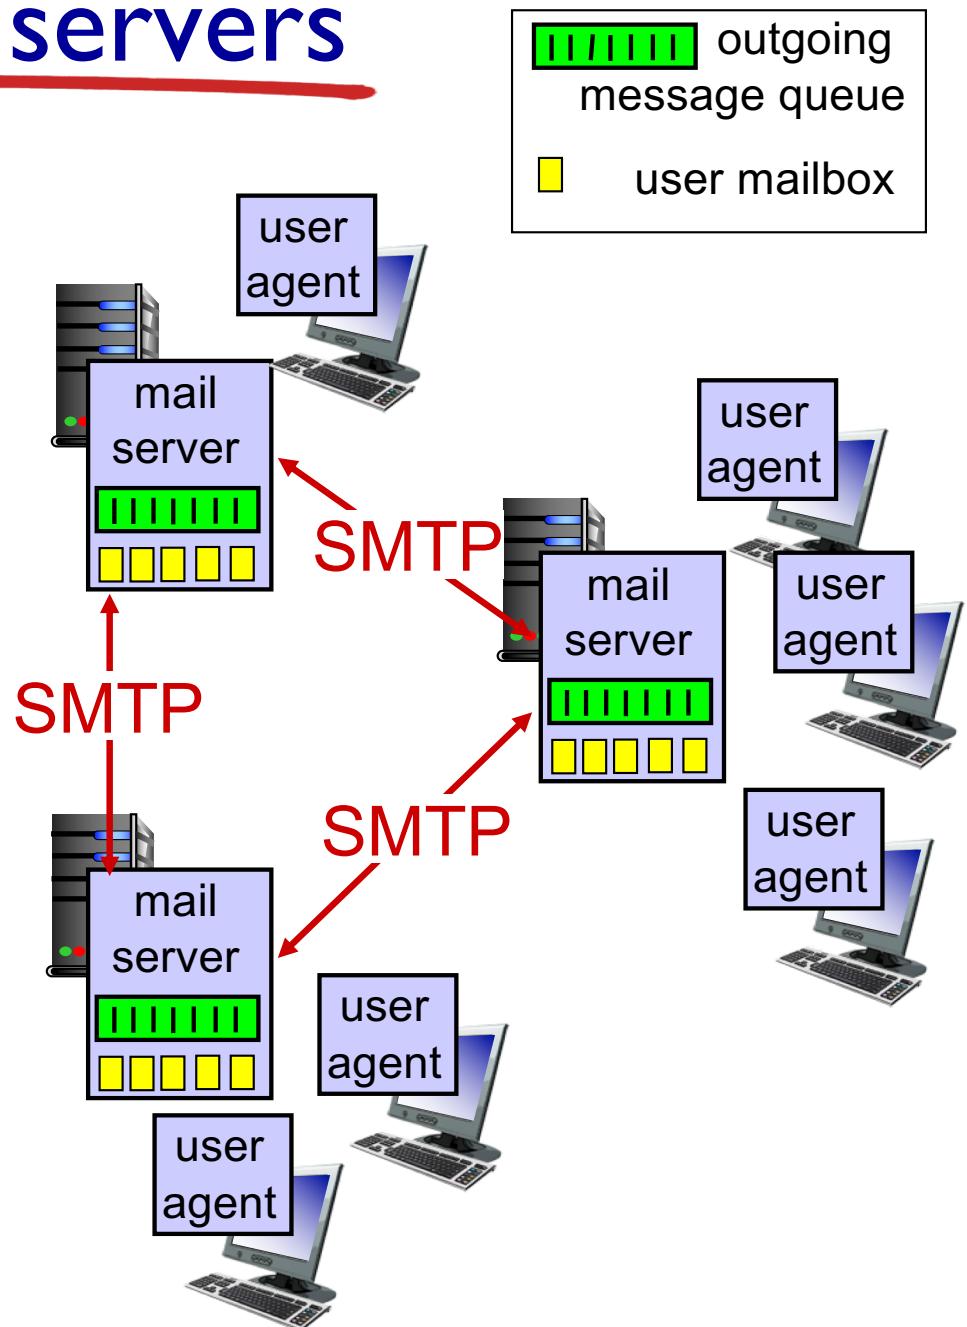
- ❖ a.k.a. “mail reader”
- ❖ composing, editing, reading mail messages
- ❖ e.g., Outlook, Thunderbird, iPhone mail client



Electronic mail: mail servers

mail servers:

- ❖ *mailbox* contains *incoming* messages for user
- ❖ *message queue* of *outgoing* (to be sent) mail messages
- ❖ *SMTP protocol* between mail servers to send email messages
 - client: sending mail server
 - “server”: receiving mail server



Electronic Mail: SMTP [RFC 2821]

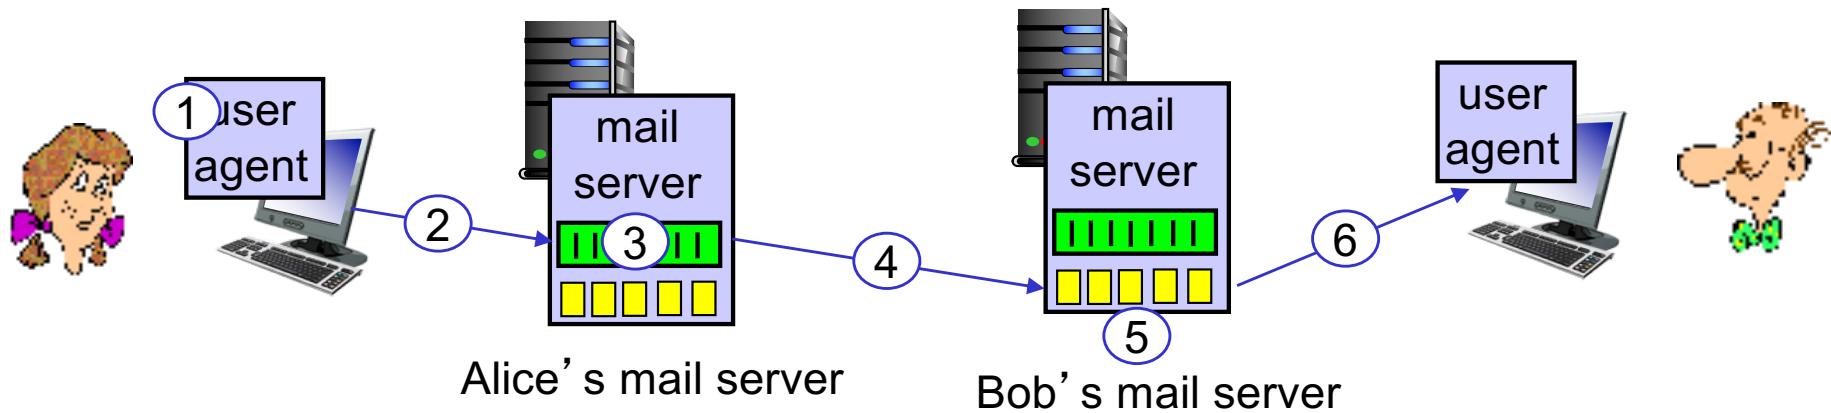
- ❖ uses **TCP** to reliably transfer email message from client to server, **port 25**
- ❖ direct transfer: sending server to receiving server
- ❖ three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure
- ❖ command/response interaction (like HTTP, FTP)
 - **commands:** ASCII text
 - **response:** status code and phrase
- ❖ messages must be in 7-bit ASCII

Sample SMTP interaction

S: 220 hamburger.edu (server's host name)
C: HELO crepes.fr (client's host name)
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr> (beginning of message)
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: . (end of message)
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection

Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



SMTP: final words

- ❖ SMTP uses **persistent connections**: If the sending mail server has several messages to send to the same receiving mail server, it can send all of the messages over the **same TCP connection**.

comparison with HTTP:

- ❖ HTTP: **pull protocol** (pull the information from the server)
SMTP: **push protocol** (sending mail server pushes the file to the receiving mail server)
- ❖ Both have ASCII command/response interaction, status codes

Mail message format

SMTP: protocol for
exchanging email msgs

RFC 822: standard for text
message format:

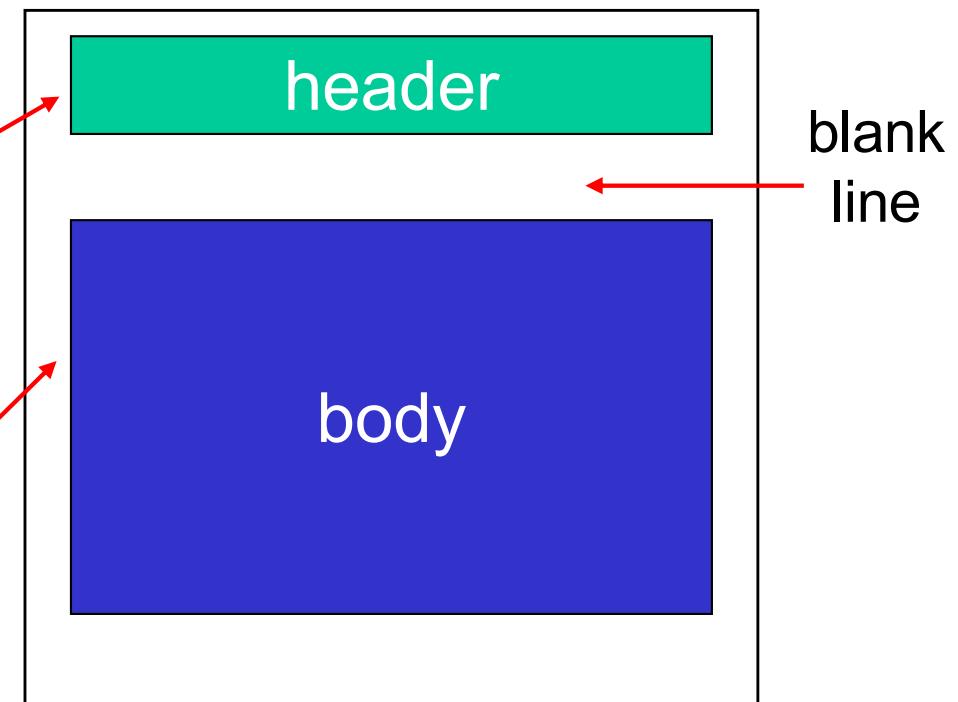
❖ **header lines**, e.g.,

- To:
- From:
- Subject:

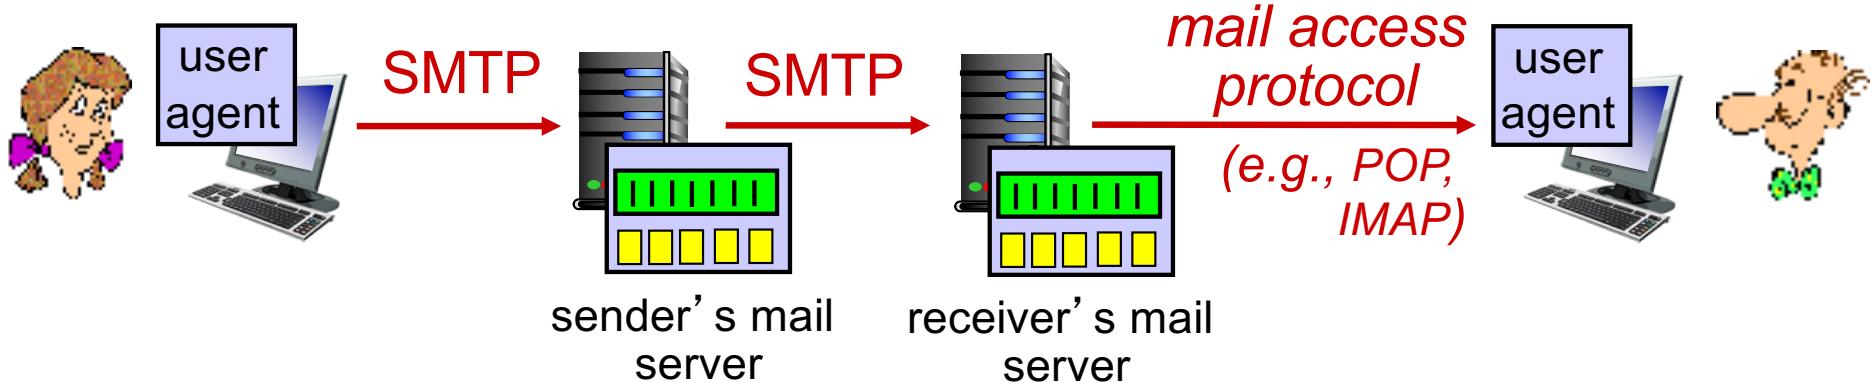
*different from SMTP MAIL
FROM, RCPT TO:
commands!*

❖ **Body**: the “message”

- ASCII characters only



Mail access protocols



- ❖ **SMTP:** delivery/storage to receiver's server
- ❖ **mail access protocol:** retrieval from server
 - **POP:** Post Office Protocol [RFC 1939]: authorization, download
 - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored msgs on server
 - **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.

POP3 protocol

- ❖ POP3 is an extremely simple mail access protocol
 - ❖ its functionality is rather limited.
- ❖ POP3 begins when the user agent (the client) opens a **TCP connection** to the mail server (the server) on **port 110**.
- ❖ With the TCP connection established, POP3 progresses through **three phases**:
 - Authorization
 - Transaction
 - Update

POP3 protocol

authorization phase

- ❖ client commands:
 - **user**: declare username
 - **pass**: password
- ❖ server responses
 - **+OK**
 - **-ERR**

transaction phase, client:

- ❖ **list**: list the size of each of the stored messages
- ❖ **retr**: retrieve message by number
- ❖ **dele**: delete
- ❖ **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 2 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

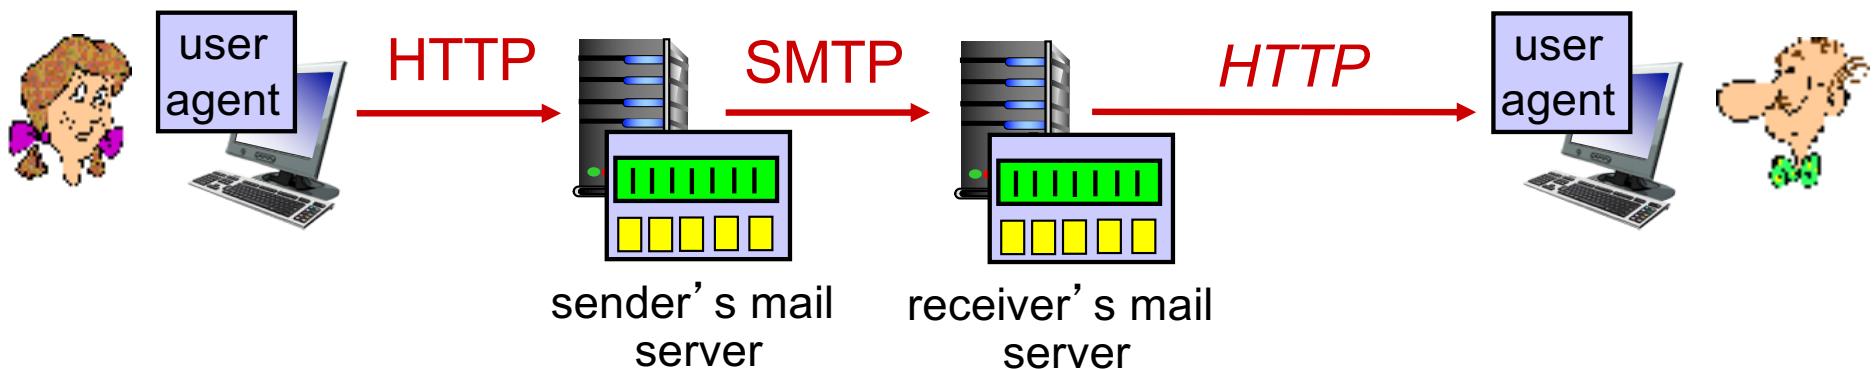
POP3 protocol

- ❖ **Update phase:** Remove messages from mailbox in this example
- ❖ Two Modes:
 - “**download and delete**”: the user agent **deletes** the messages on the mail server after downloading them
 - “**download and keep**”: the user agent **leaves** the messages on the mail server after downloading them
- ❖ POP3 is **stateless** across sessions.

IMAP protocol

- ❖ keeps all messages in one place: at server
- ❖ allows user to organize messages in folders
- ❖ keeps user **state** across sessions:
 - names of folders and mappings between message IDs and folder name

Web-Based E-mail



- ❖ **User agent:** Web browser
- ❖ **User communicates with remote mailbox via HTTP:**
 - **Send e-mail:** e-mail message is sent from browser to mail server over HTTP (not SMTP)
 - **Mail server message exchange:** Still using SMTP
 - **Access message in mailbox:** e-mail message is sent from mail server to browser using the HTTP (not POP3 or IMAP)

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Socket programming with UDP and TCP

DNS: domain name system

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System (DNS):

- ❖ *distributed database* implemented in hierarchy of many *name servers*
- ❖ *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)
 - note: core Internet function, implemented as application-layer protocol
 - complexity at network’s “edge”

DNS: services

- ❖ **hostname to IP address translation**
 - **Hostname**: mnemonic identifier for a host
 - E.g., www.yahoo.com
- ❖ **host aliasing**
 - E.g., **canonical hostname**: relay1.west-coast.enterprise.com
 - Two **alias**: enterprise.com and www.enterprise.com (more mnemonic)
- ❖ **mail server aliasing**
 - E.g., E-mail address: bob@hotmail.com
 - Alias hostname: hotmail.com
 - Canonical hostname: relay1.west-coast.hotmail.com
 - Note: a company's mail server and Web server can have identical (aliased) hostnames, e.g., enterprise.com.
- ❖ **load distribution**
 - replicated Web servers: many IP addresses correspond to one name

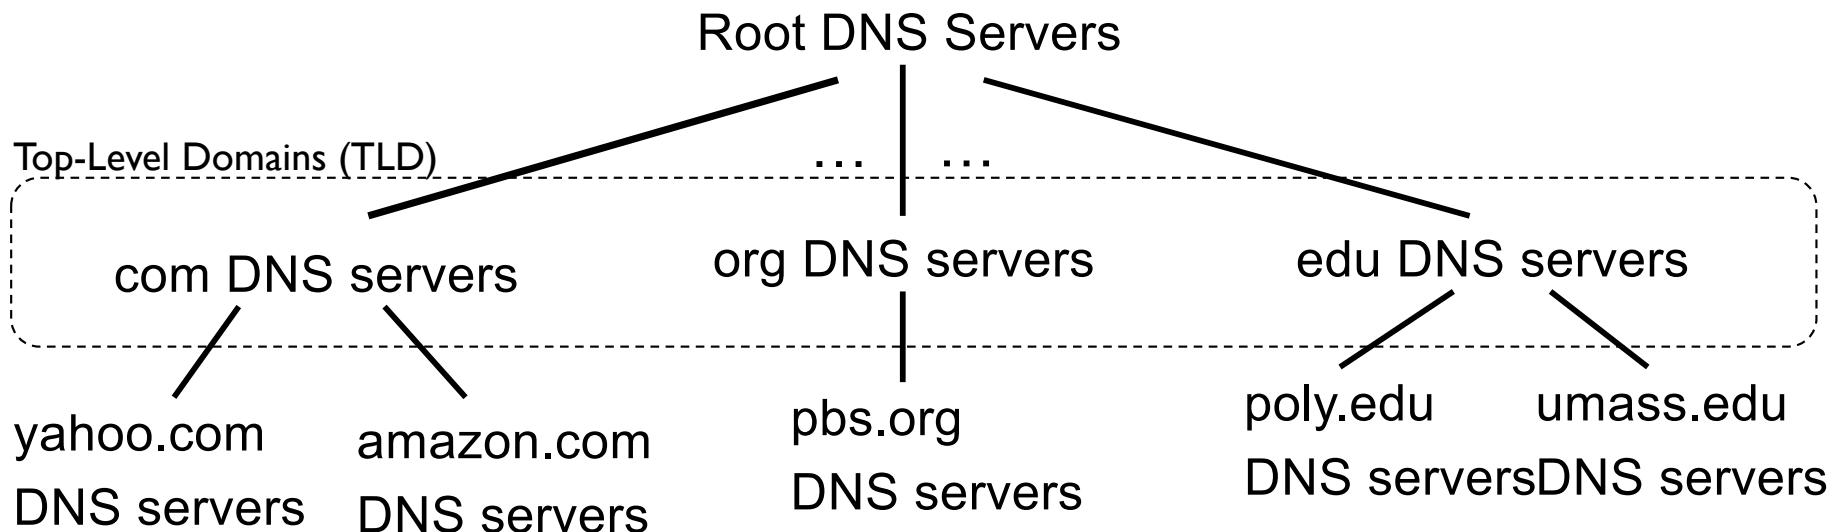
DNS structure

why not centralize DNS?

- ❖ single point of failure
- ❖ traffic volume
- ❖ distant centralized database
- ❖ maintenance

A: doesn't scale!

DNS: a distributed, hierarchical database



Principle

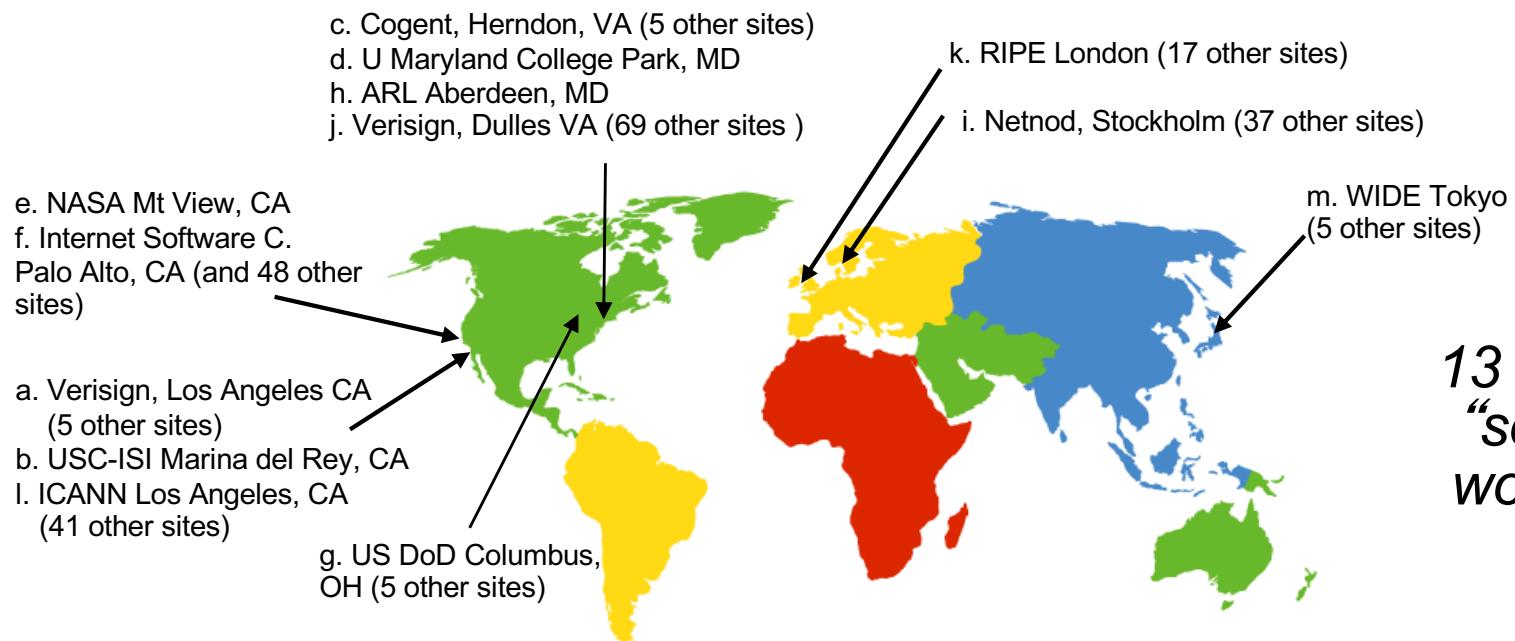
- ❖ parent server has IP-name records for children

To resolve IP for www.amazon.com; 1st approx:

- ❖ queries root server to find .com DNS server
- ❖ client queries .com DNS server to get amazon.com DNS server
- ❖ client queries amazon.com DNS server to get IP address for www.amazon.com

DNS: root name servers

- ❖ contacted by local name server that can not resolve name
- ❖ **root name server:**
 - contacts authoritative name server if name mapping not known
→ gets mapping
 - returns mapping to local name server



*13 root name
“servers”
worldwide*

TLD, authoritative servers

Top-level domain (TLD) servers:

- responsible for **com, org, net, edu, aero, jobs, museums,** and all top-level country domains, e.g.: **uk, fr, ca, jp**
- further delegates query to authoritative DNS servers

Authoritative DNS servers:

- **organization's own DNS server(s)**, providing authoritative hostname to IP mappings for organization's named hosts (e.g., *.cuhk.edu.hk), known as a DNS zone
- can be maintained by organization or service provider
- can further sub-divide into **smaller DNS zones** each with their own authoritative DNS server (e.g., *.ie.cuhk.edu.hk)

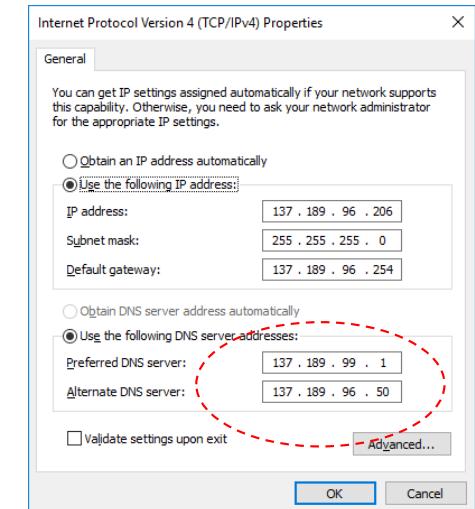
Local DNS name server

- ❖ does not strictly belong to hierarchy
 - ❖ each ISP (residential ISP, company, university) has one
 - also called “**default name server**”
 - ❖ when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

```
C:\>ipconfig /all

Ethernet adapter Ethernet:

  Connection-specific DNS Suffix  . :
  Description . . . . . : Intel(R) Ethernet Connection I217-LM
  Physical Address. . . . . : 64-00-6A-5F-70-A3
  DHCP Enabled. . . . . : No
  Autoconfiguration Enabled . . . . : Yes
  IPv4 Address. . . . . : 137.189.96.206(Preferred)
  Subnet Mask . . . . . : 255.255.255.0
  Default Gateway . . . . . : 137.189.96.254
  DHCPv6 IAID . . . . . : 56885354
  DHCPv6 Client DUID. . . . . : 00-01-00-01-1D-96-3D-F9-64-00-6A-5F-
  70-A3
  DNS Servers . . . . . : 137.189.99.1
                           137.189.96.50
```

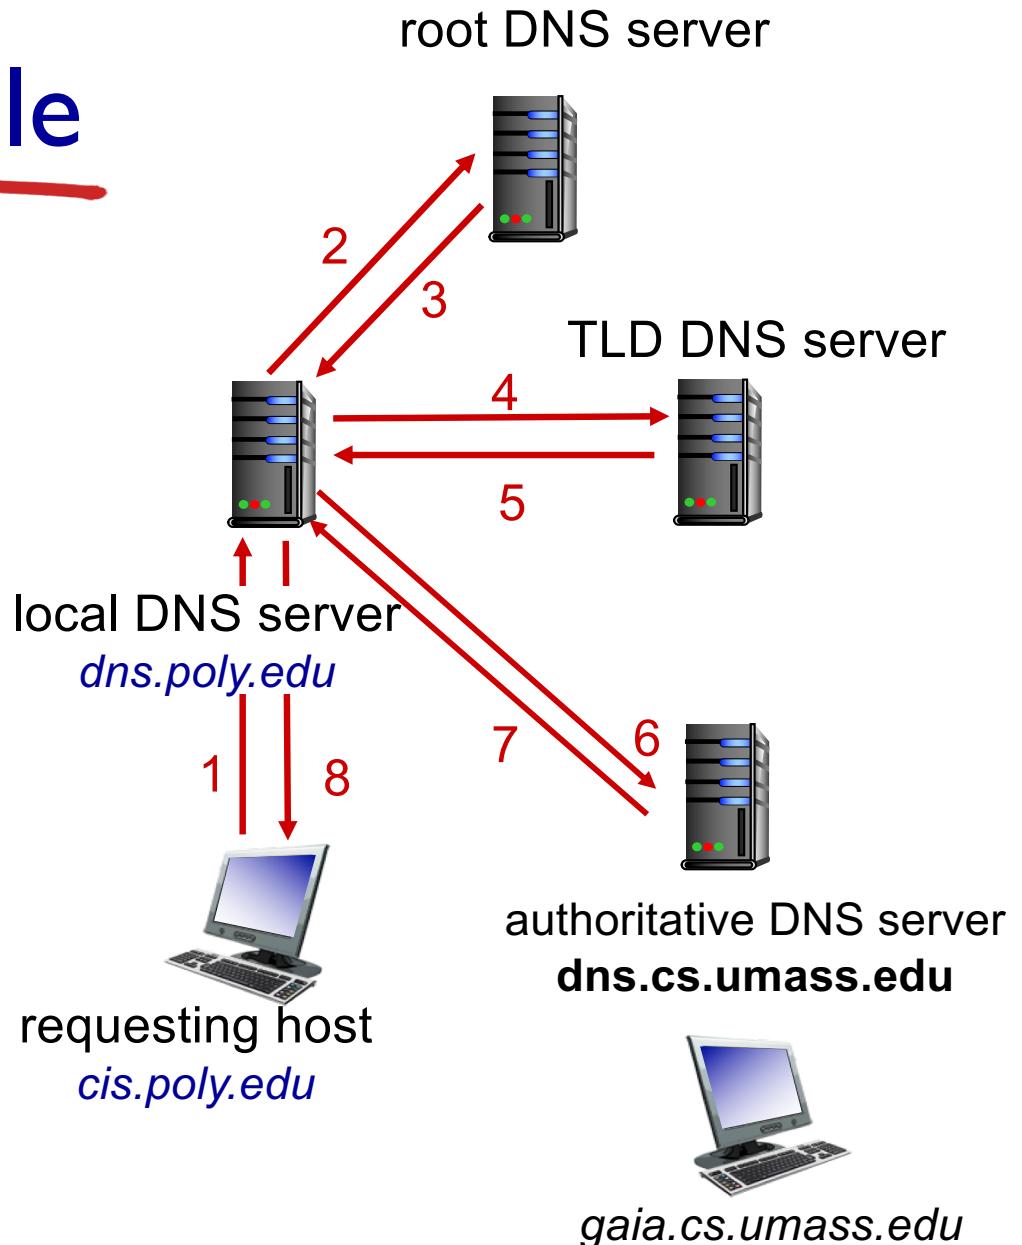


DNS name resolution example

- ❖ *Example:* host at `cis.poly.edu` wants IP address for `gaia.cs.umass.edu`

Recursive query:

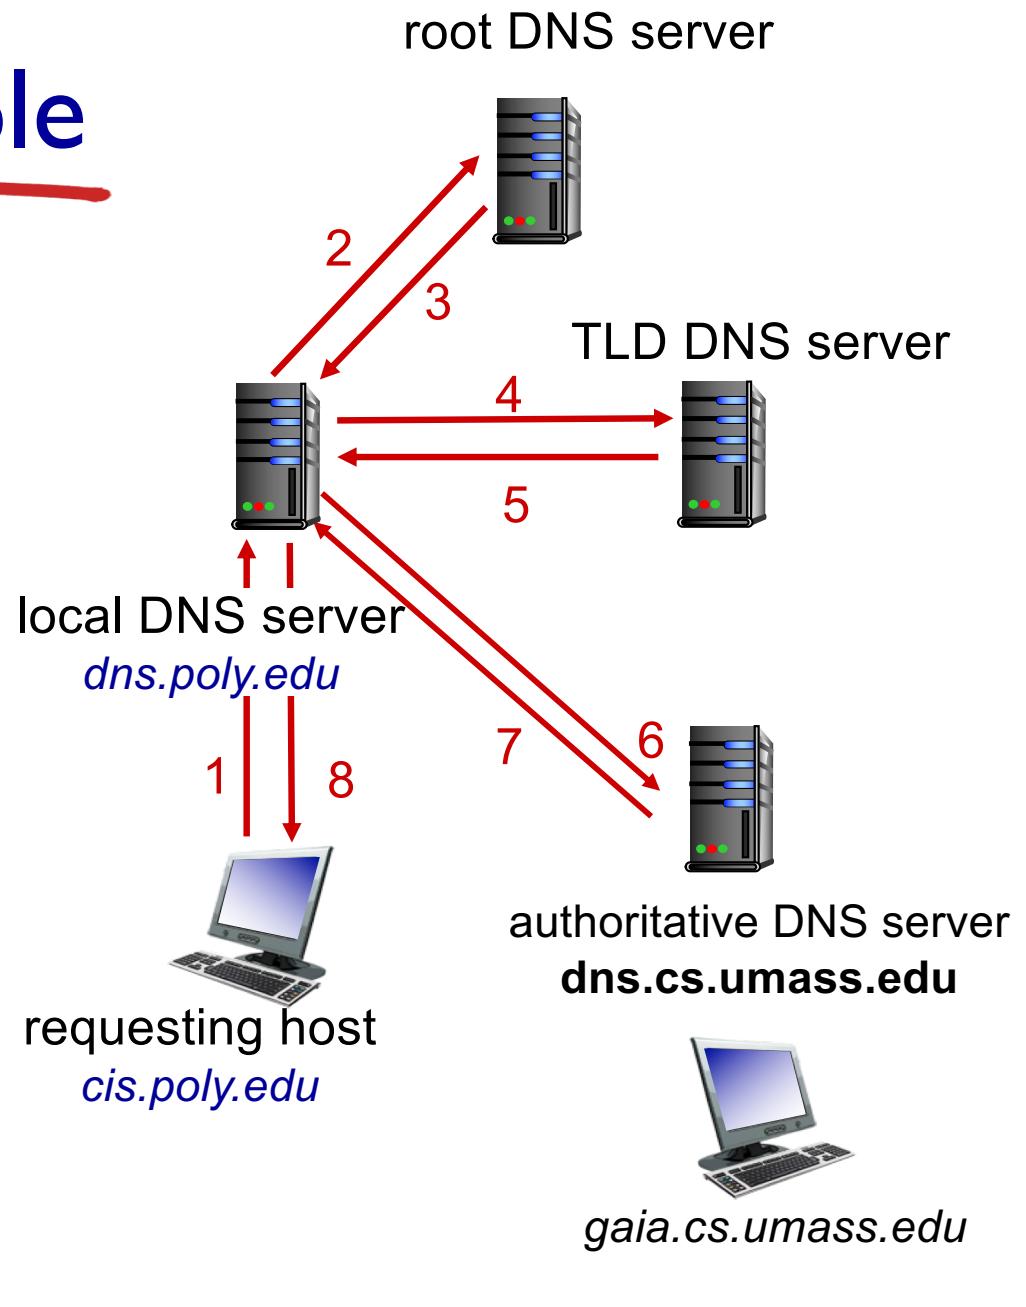
- ❖ puts burden of name resolution on contacted name server
- ❖ E.g., (1,8)



DNS name resolution example

Iterative query:

- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”
- ❖ E.g., (2,3), (4,5), (6,7)

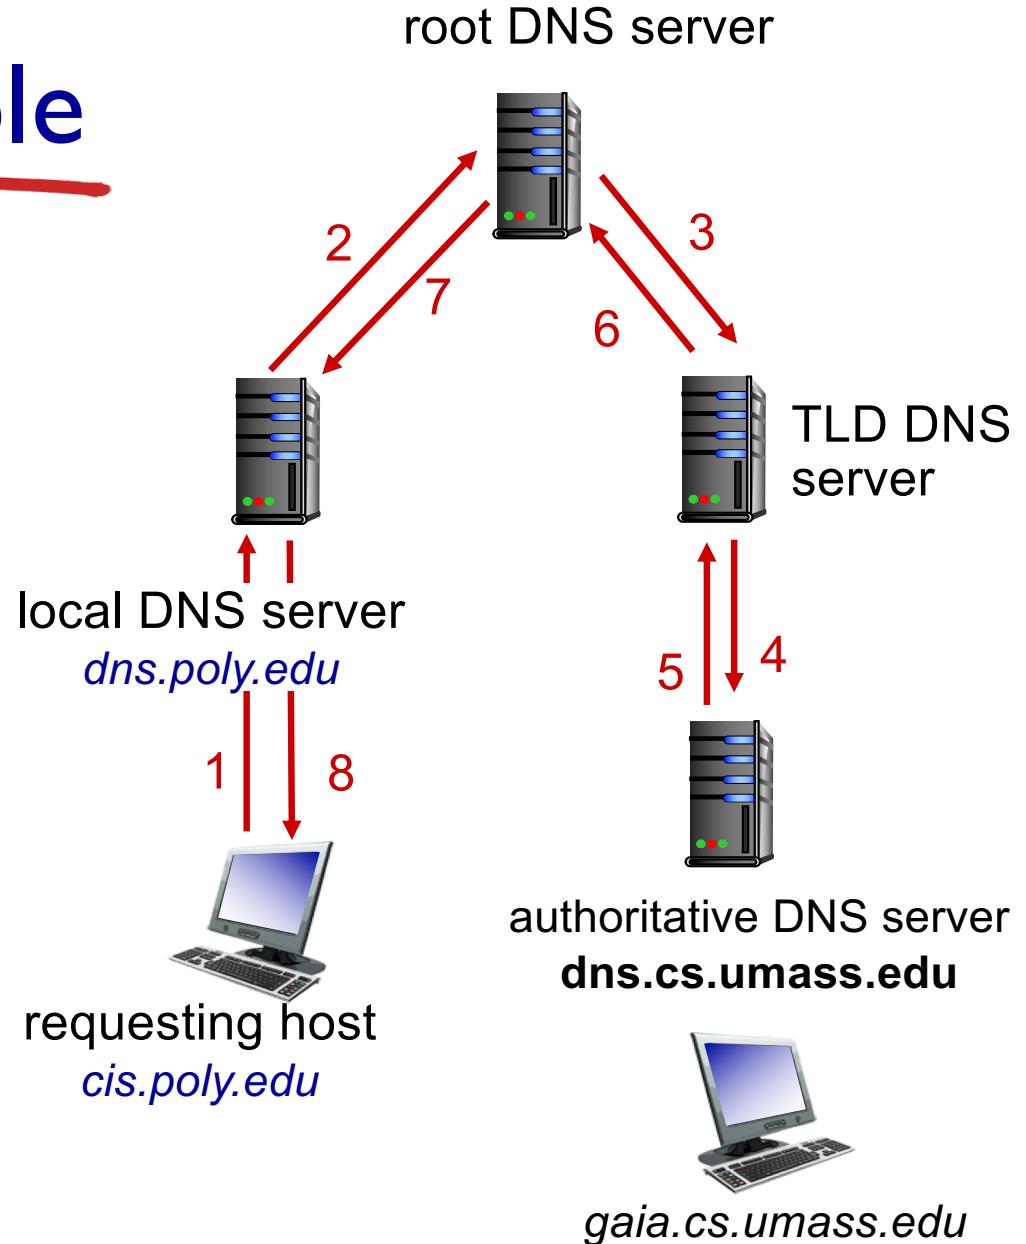


Typical query pattern:

- ❖ query from the requesting host to the local DNS server is recursive
- ❖ remaining queries are iterative.

DNS name resolution example

- ❖ Another possible query pattern: All the queries are recursive
- ❖ heavy load at upper levels of hierarchy



DNS caching

- ❖ once (any) name server learns mapping, it *caches* it
 - Adv: improve delay performance and to reduce the number of DNS messages ricocheting around
 - Query for same hostname: DNS server can provide desired IP address, even if it is not authoritative
- ❖ cached entries may be *out-of-date* (best effort name-to-address translation!)
 - cache entries removed after **time to live (TTL)**
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- ❖ update/notify mechanisms proposed IETF standard
 - RFC 2136

DNS records

DNS: distributed database storing **resource records (RR)**

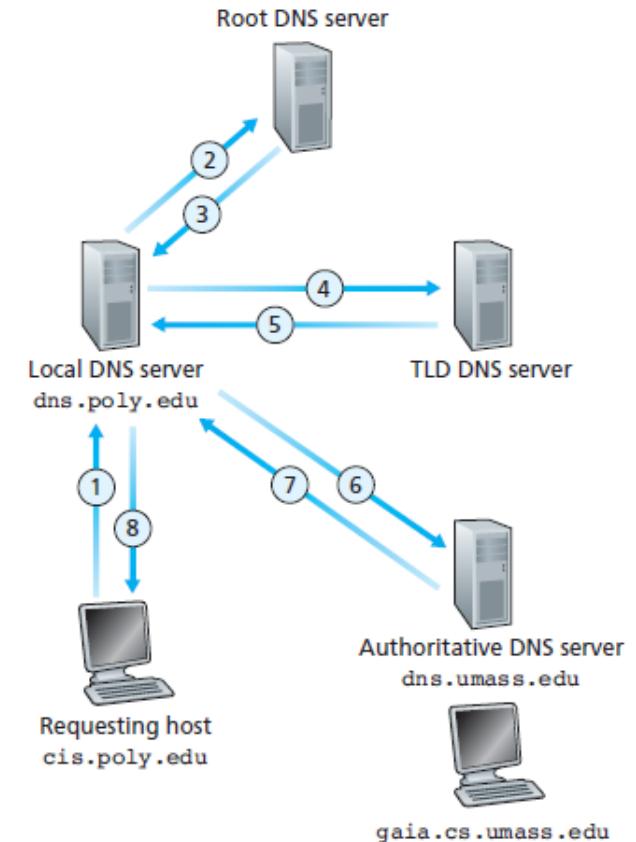
RR format: `(name, value, type, ttl)`

type=A

- **name** is hostname (e.g., `relay1.bar.foo.com`)
- **value** is IP address

type=NS

- **name** is domain (e.g., `foo.com`)
- **value** is hostname of authoritative name server for this domain (e.g., `dns.foo.com`)



DNS records (cont'd)

RR format: (**name**, **value**, **type**, **ttl**)

type=CNAME

- **name** is alias name
- **value** is canonical (the real) name
- E.g., (foo.com, relay1.bar.foo.com, CNAME, ttl)

type=MX

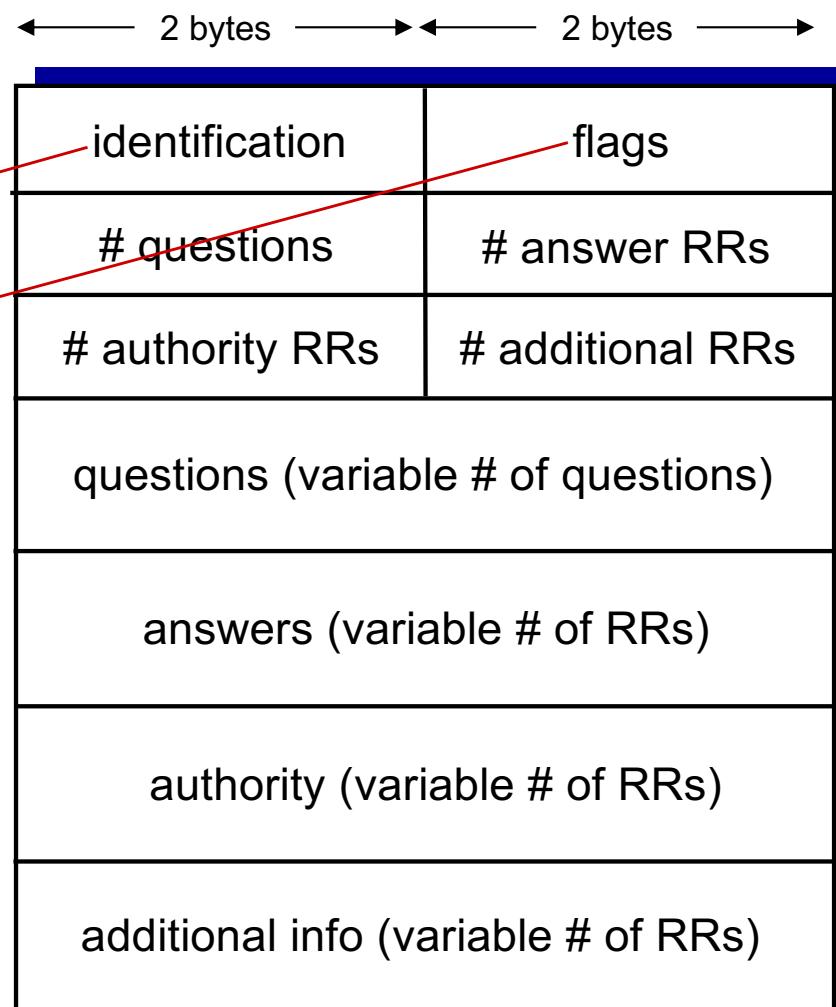
- **value** is name of mail server associated with **name**
- E.g., (foo.com, mail.bar.foo.com, MX, ttl)

DNS protocol, messages

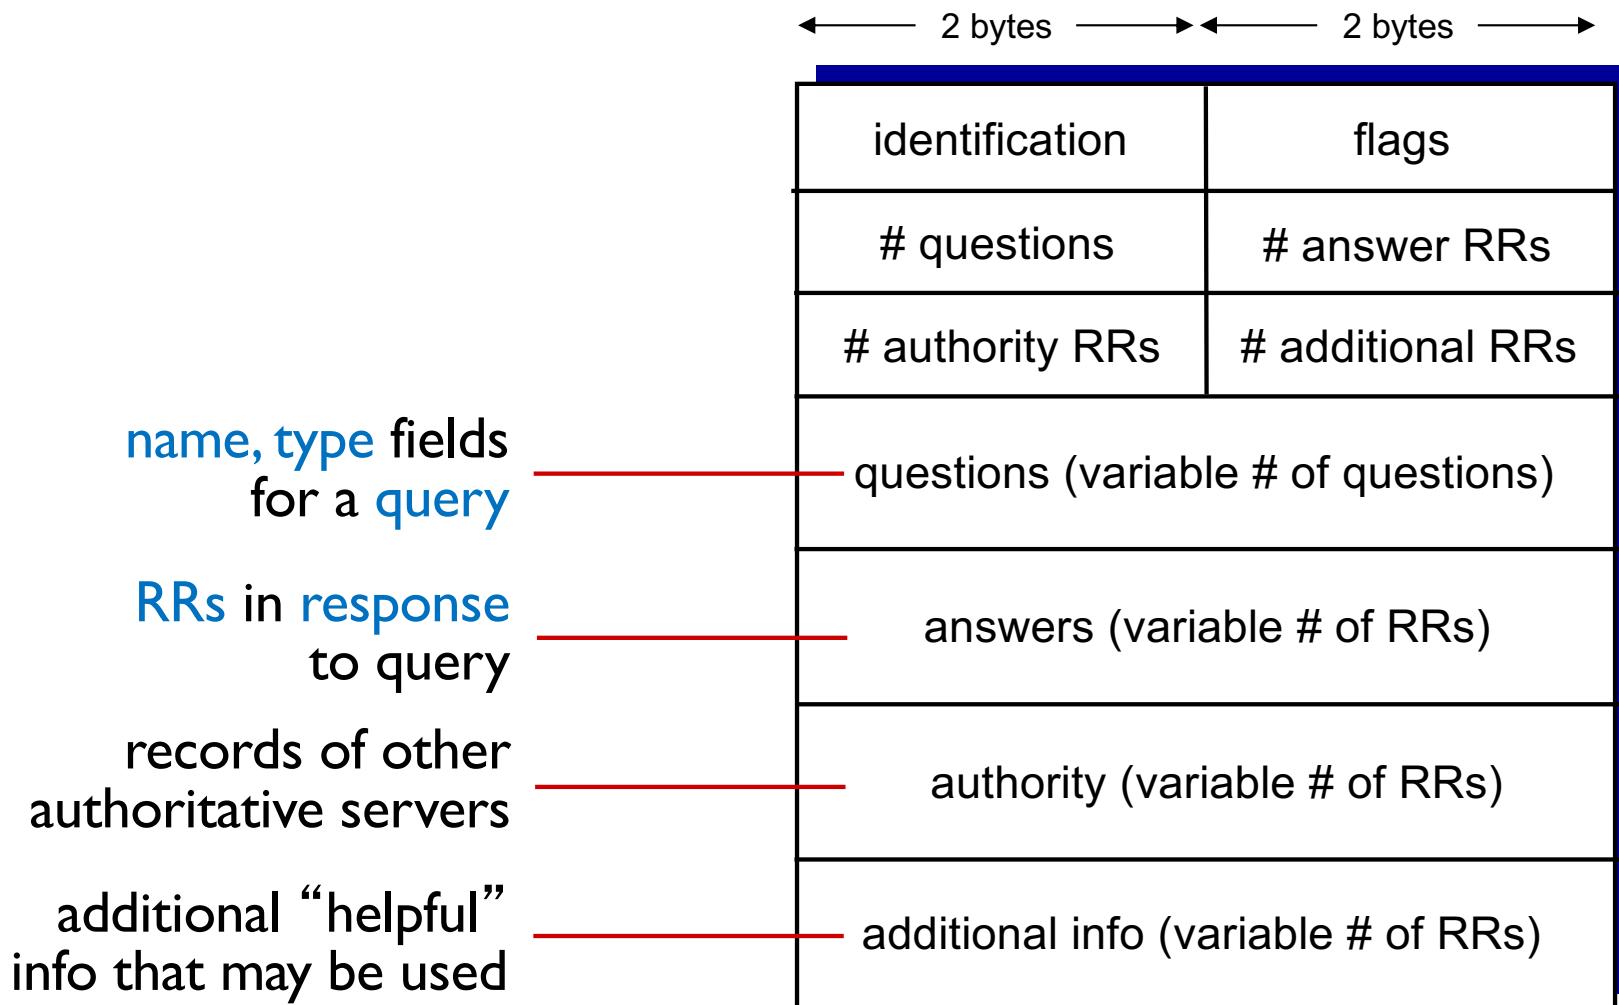
- ❖ *query* and *reply* messages, both with same *message format*

msg header

- ❖ **identification:** 16 bit # for query,
reply to query uses same #
- ❖ **flags:**
 - query or reply
 - recursion desired by client
(if DNS server does not
have the record)
 - recursion supported
 - reply is authoritative



DNS protocol, messages



Inserting records into DNS

- ❖ example: new startup “Network Utopia”
- ❖ register name networkuptopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of **authoritative server** (primary and secondary)
 - registrar inserts two RRs into .com TLD server:
(`networkutopia.com`, `dns1.networkutopia.com`, NS)
(`dns1.networkutopia.com`, 212.212.212.1, A)
- ❖ create RRs at your **authoritative server**
 - type A record for Web server `www.networkuptopia.com`

Attacking DNS

DDoS attacks

- ❖ Bombard root servers with traffic
 - Not successful to date
 - Traffic Filtering
 - Local DNS servers cache IPs of TLD servers, allowing root server bypass
- ❖ Bombard TLD servers
 - Potentially more dangerous

Redirect attacks

- ❖ Man-in-middle
 - Intercept queries
- ❖ DNS poisoning
 - Send bogus replies to DNS server, which caches

Exploit DNS for DDoS

- ❖ Send queries with spoofed source address: target IP
- ❖ Requires amplification

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

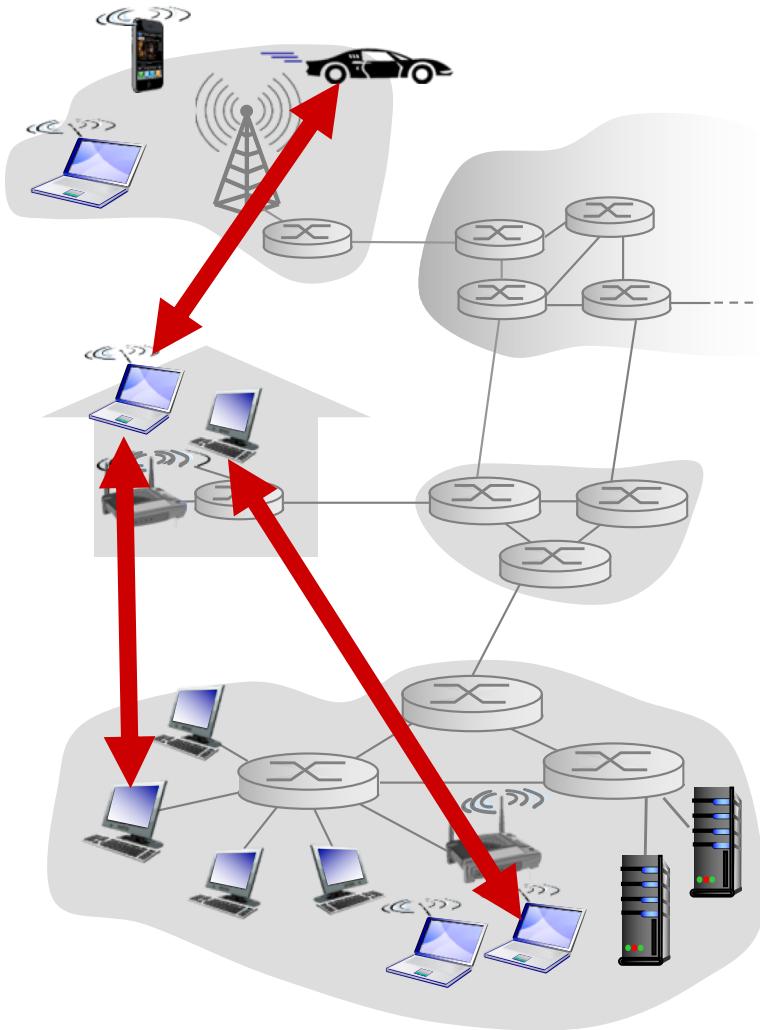
2.7 Socket programming with UDP and TCP

Pure P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers are intermittently connected and change IP addresses

examples:

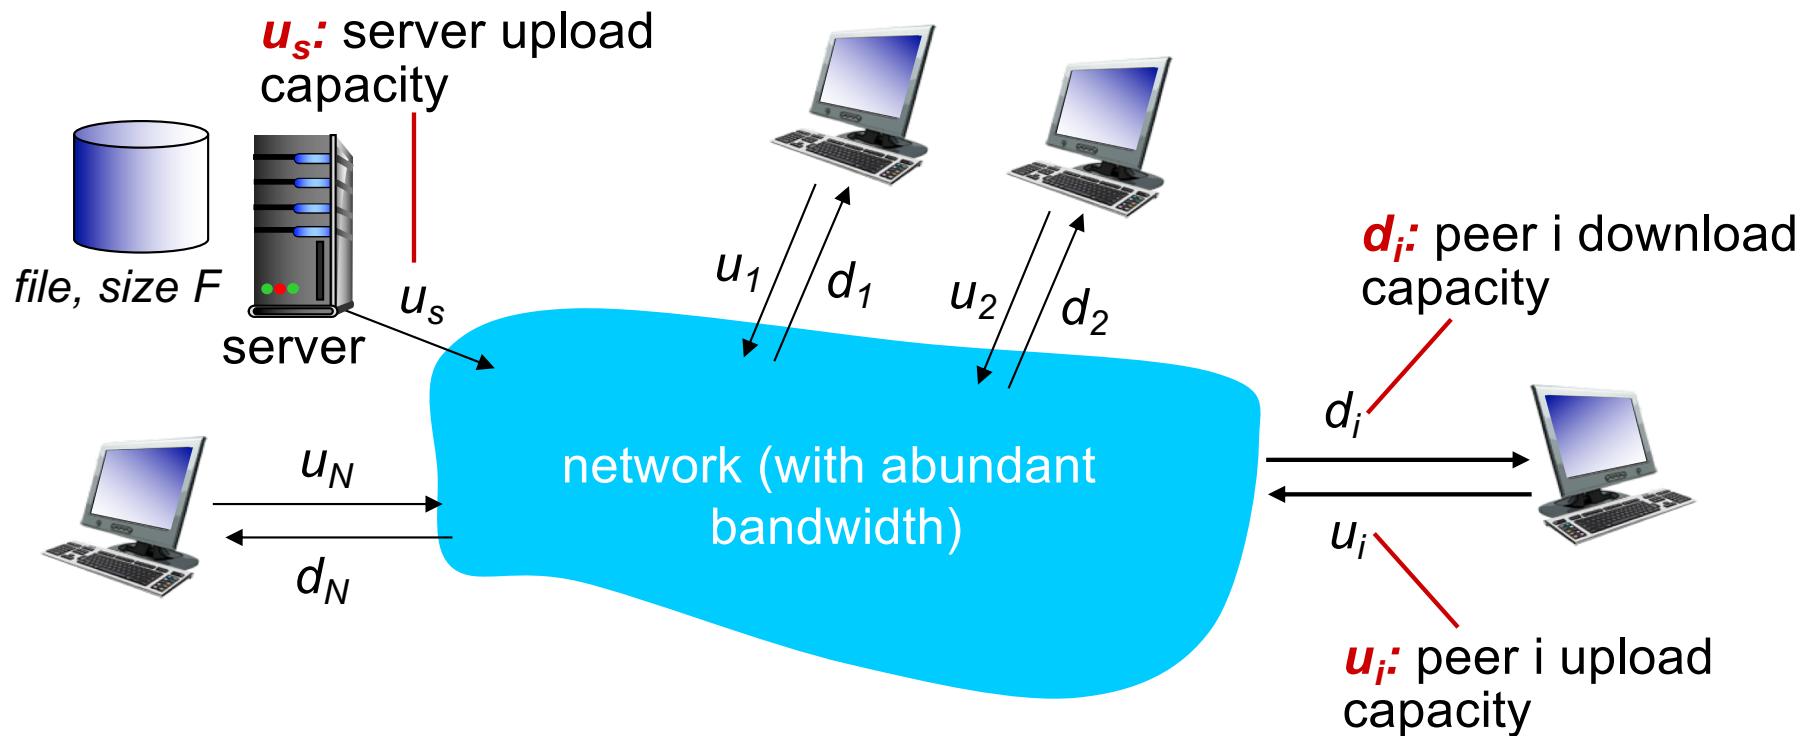
- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



File distribution: client-server vs P2P

Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource

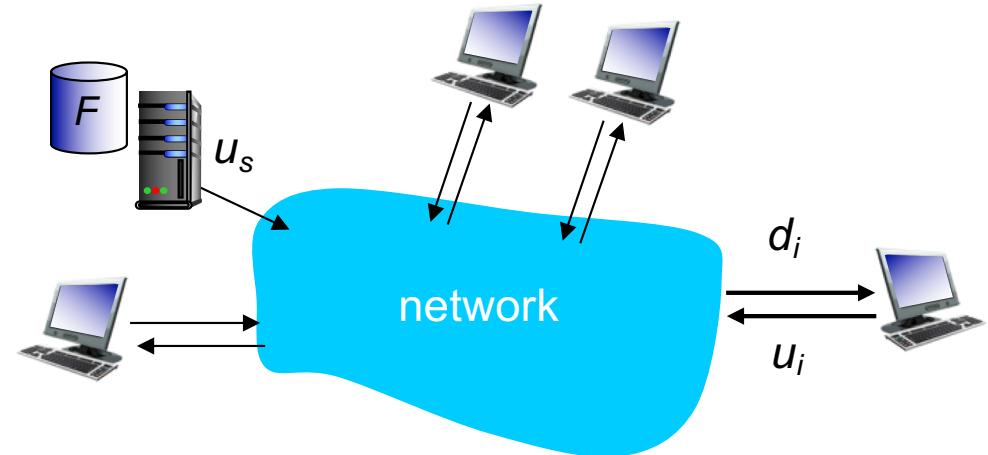


File distribution time: client-server

- ❖ **server transmission:** must sequentially send (upload) N file copies:

- time to send one copy: F/u_s
- time to send N copies: NF/u_s

- ❖ **client:** each client must download file copy
 - d_{\min} = min client download rate
 - min client download time: F/d_{\min}



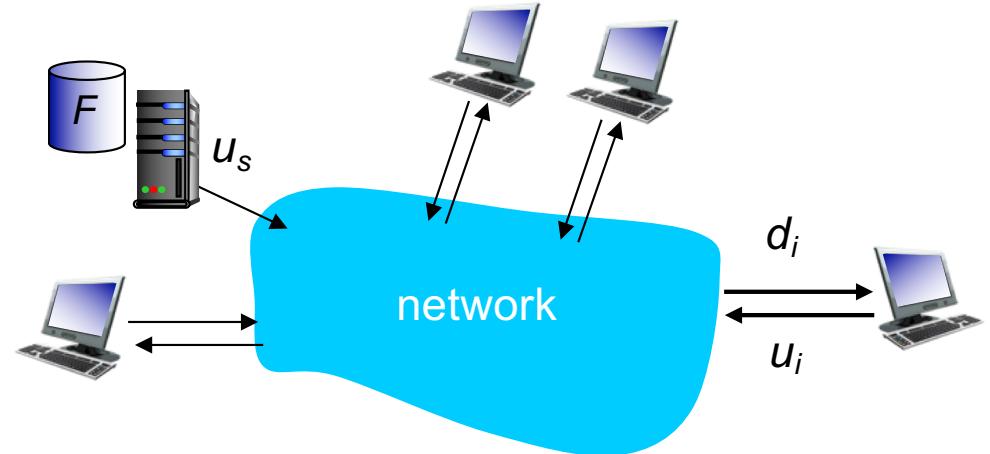
*time to distribute F
to N clients using
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N

File distribution time: P2P

- ❖ **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- ❖ **client:** each client must download file copy
 - min client download time: F/d_{\min}
- ❖ **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$



time to distribute F

to N clients using
P2P approach

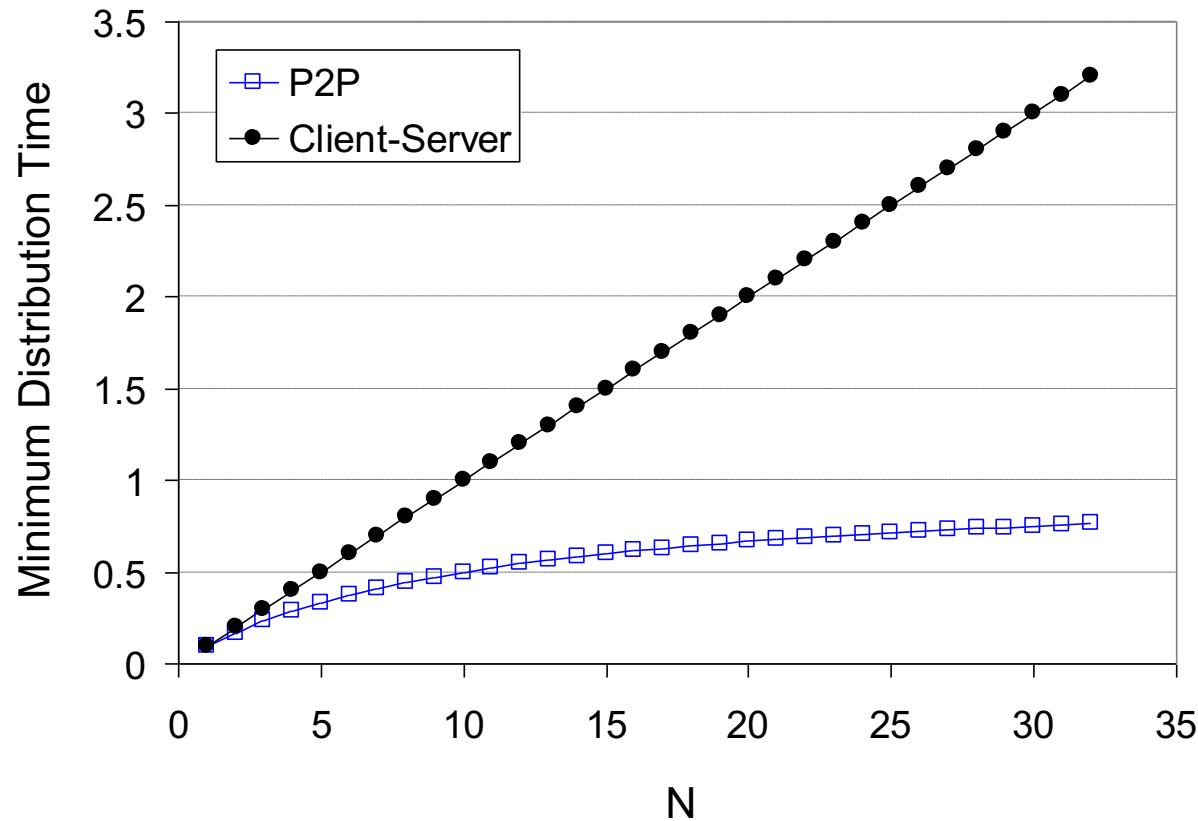
$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



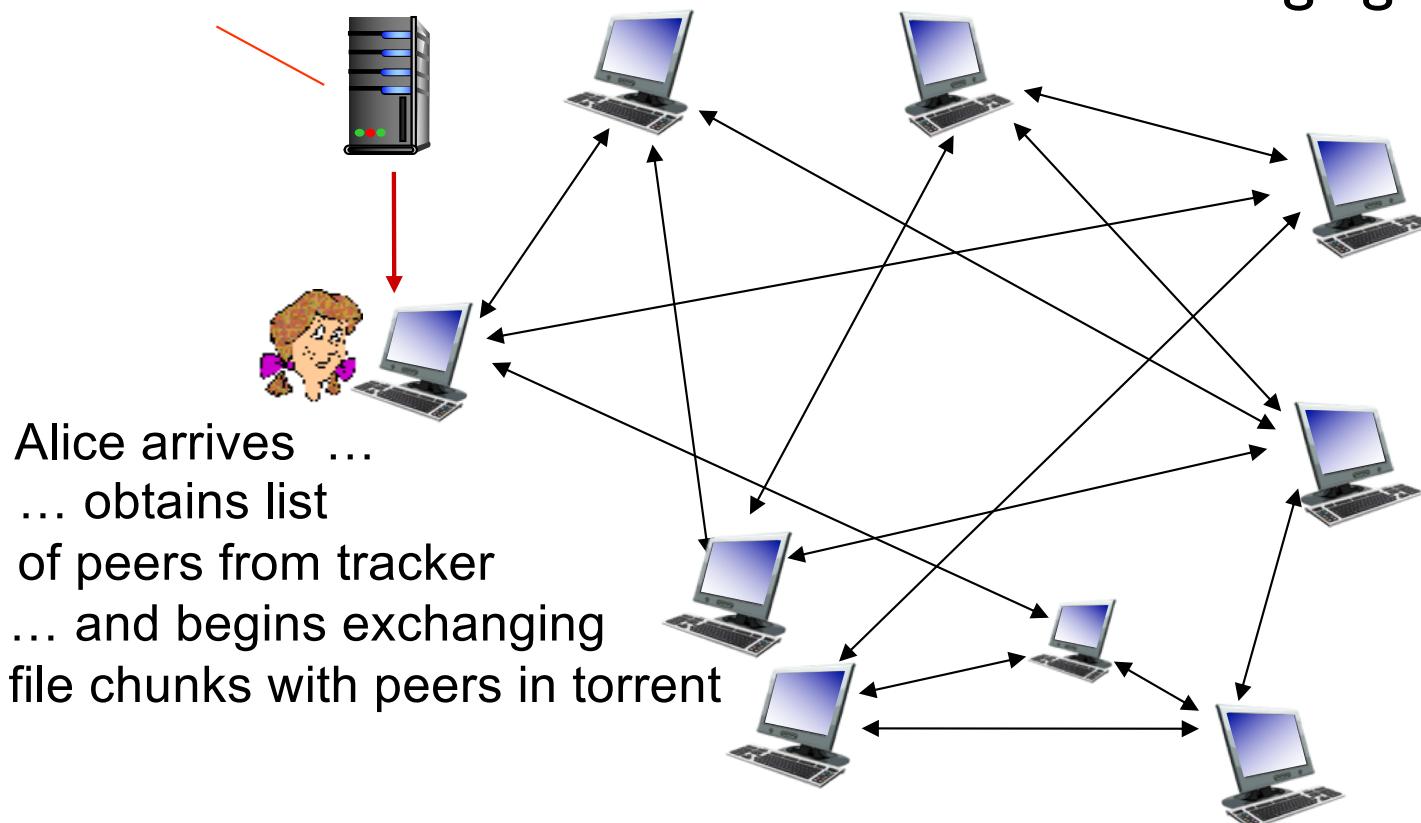
- **Scalability:** Minimal distribution time of P2P architecture is always less than that of the client-server architecture.
- **Reason:** Peers serve as **redistributors** as well as **consumers** of bits.

P2P file distribution: BitTorrent

- ❖ file divided into 256Kb **chunks**
- ❖ peers in torrent send/receive file chunks

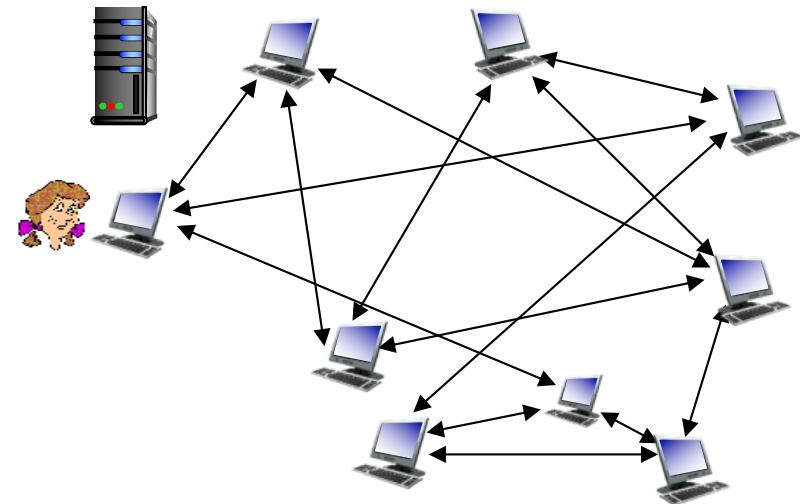
tracker: tracks peers
participating in torrent

torrent: group of peers
exchanging chunks of a file



P2P file distribution: BitTorrent

- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- ❖ while downloading, peer uploads chunks to other peers
- ❖ peer may change peers with whom it exchanges chunks
- ❖ *churn*: peers may come and go
 - ❖ once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

requesting chunks:

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, Alice asks each peer for **list of chunks** that they have
- ❖ Alice requests missing chunks from peers, **rarest first**

=> rarest chunks get more quickly redistributed

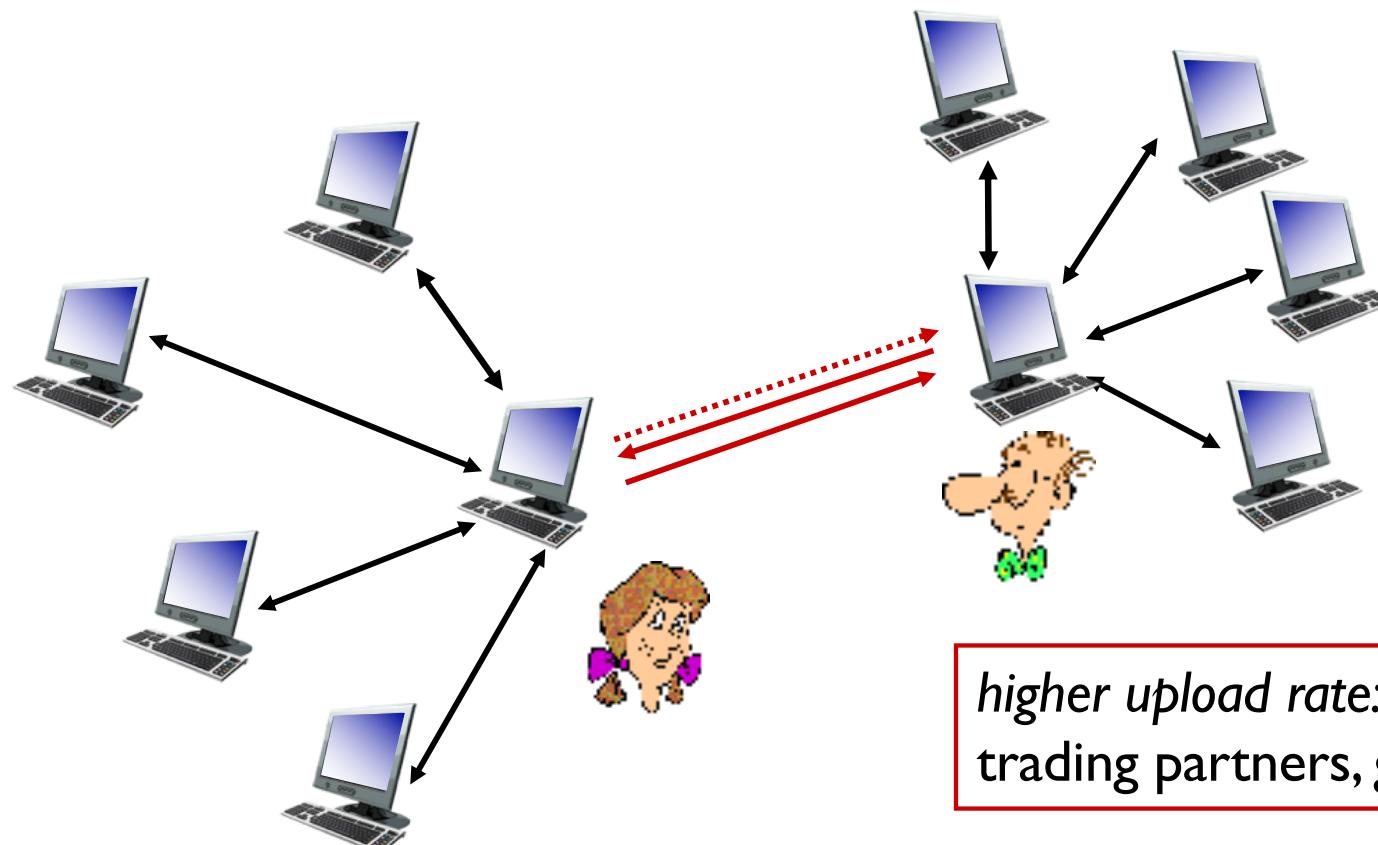
=> (roughly) equalize the numbers of copies of each chunk in the torrent

sending chunks: tit-for-tat

- ❖ Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Chapter 2: summary

our study of network apps now complete!

- ❖ application architectures
 - client-server
 - P2P
- ❖ application service requirements:
 - reliability, bandwidth, delay
- ❖ Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- ❖ specific protocols:
 - HTTP
 - FTP
 - SMTP, POP, IMAP
 - DNS
 - P2P: BitTorrent
- ❖ socket programming: TCP, UDP sockets

Chapter 2: summary

most importantly: learned about protocols!

- ❖ typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- ❖ message formats:
 - headers: fields giving info about data
 - data: info being communicated

important themes:

- ❖ control vs. data msgs
 - in-band, out-of-band
- ❖ centralized vs. decentralized
- ❖ stateless vs. stateful
- ❖ reliable vs. unreliable msg transfer
- ❖ “complexity at network edge”